

# Image Segmentation



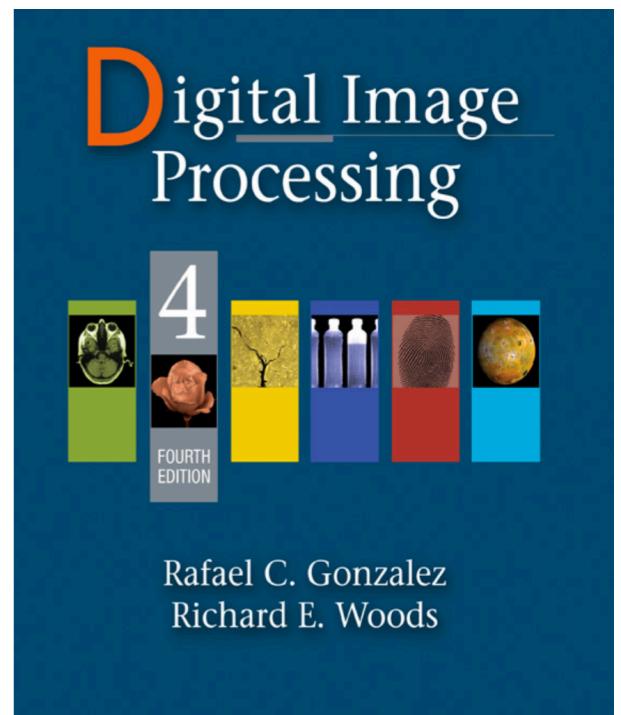
# Image Segmentation

Image segmentation is the operation of partitioning an image into a collection of sets of pixels.

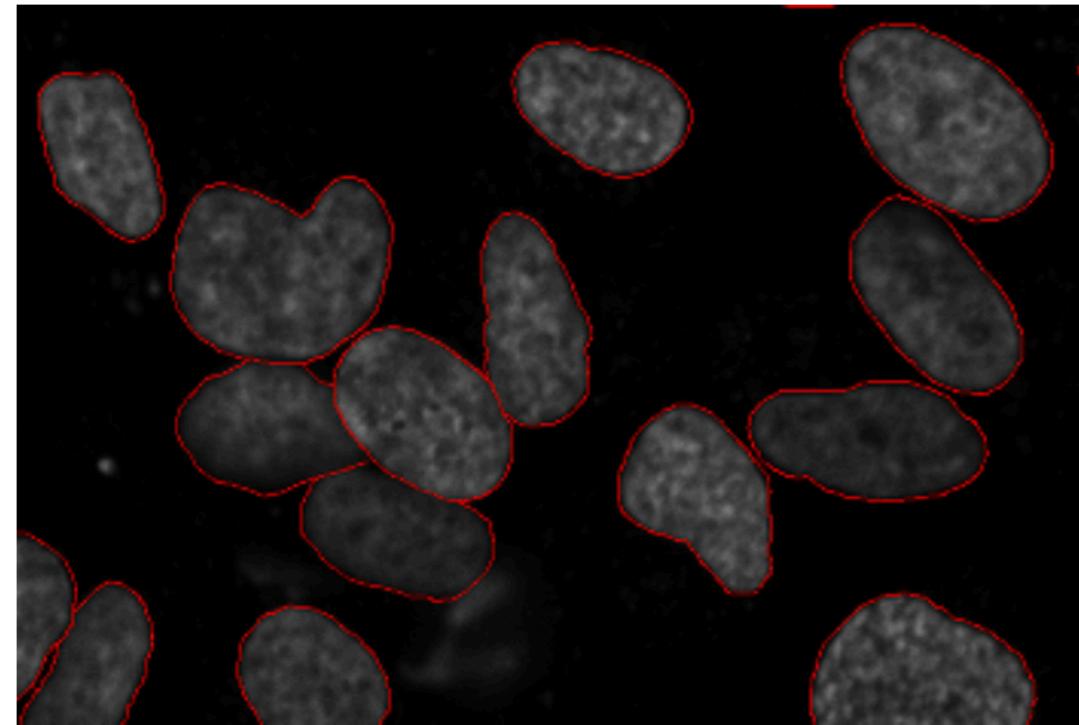
1. into regions, which usually cover the entire nimage
2. into linear structures, such as
  - line segments
  - curve segments
3. into 2D shapes, such as
  - circles
  - ellipses
  - ribbons (long, symmetric regions)

- a.  $\bigcup_{i=1}^n R_i = R$ .
- b.  $R_i$  is a connected set, for  $i = 0, 1, 2, \dots, n$ .
- c.  $R_i \cap R_j = \emptyset$  for all  $i$  and  $j$ ,  $i \neq j$ .
- d.  $Q(R_i) = \text{TRUE}$  for  $i = 0, 1, 2, \dots, n$ .
- e.  $Q(R_i \cup R_j) = \text{FALSE}$  for any adjacent regions  $R_i$  and  $R_j$

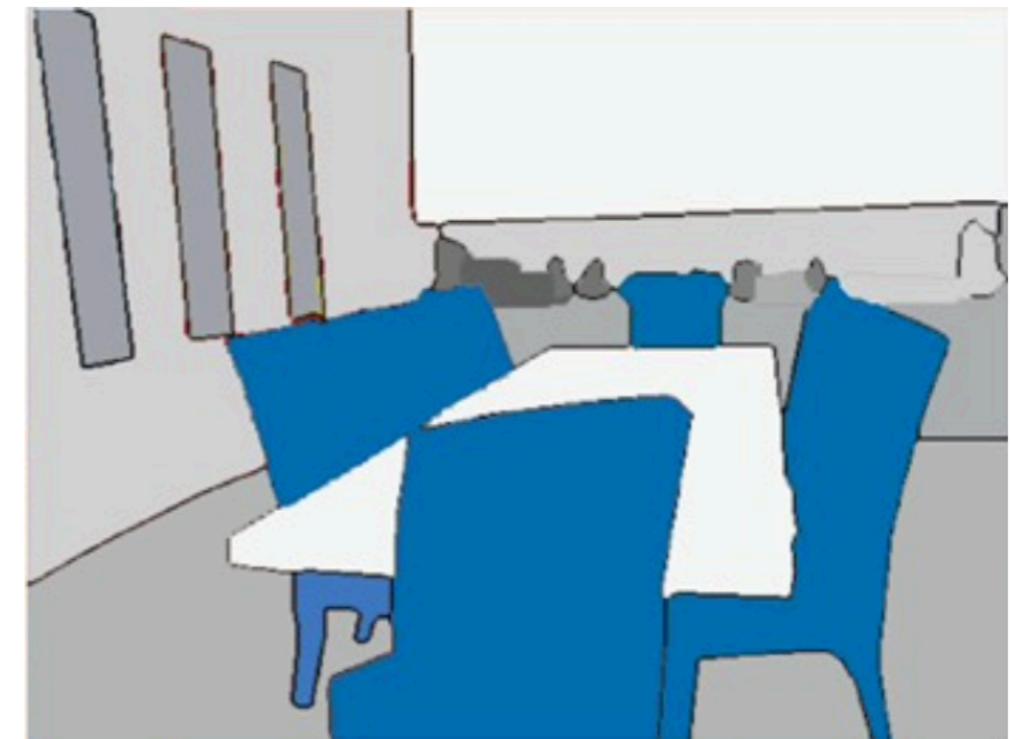
Condition (d) deals with the properties that must be satisfied by the pixels in a segmented region for example if all pixels in have the same intensity.



Separation foreground-background



Semantic segmentation



Instance segmentation



Instance segmentation is an enhanced type of object detection that generates a segmentation map for each detected instance of an object

Semantic segmentation is the task of clustering parts of an image together which belong to the same object class

# Thresholding and Binarization



Global and Local Thresholding

the threshold value  
 $T(x,y)$   
is a mean of the  
blockSize $\times$ blockSize  
neighborhood of  
( $x,y$ )  
minus C

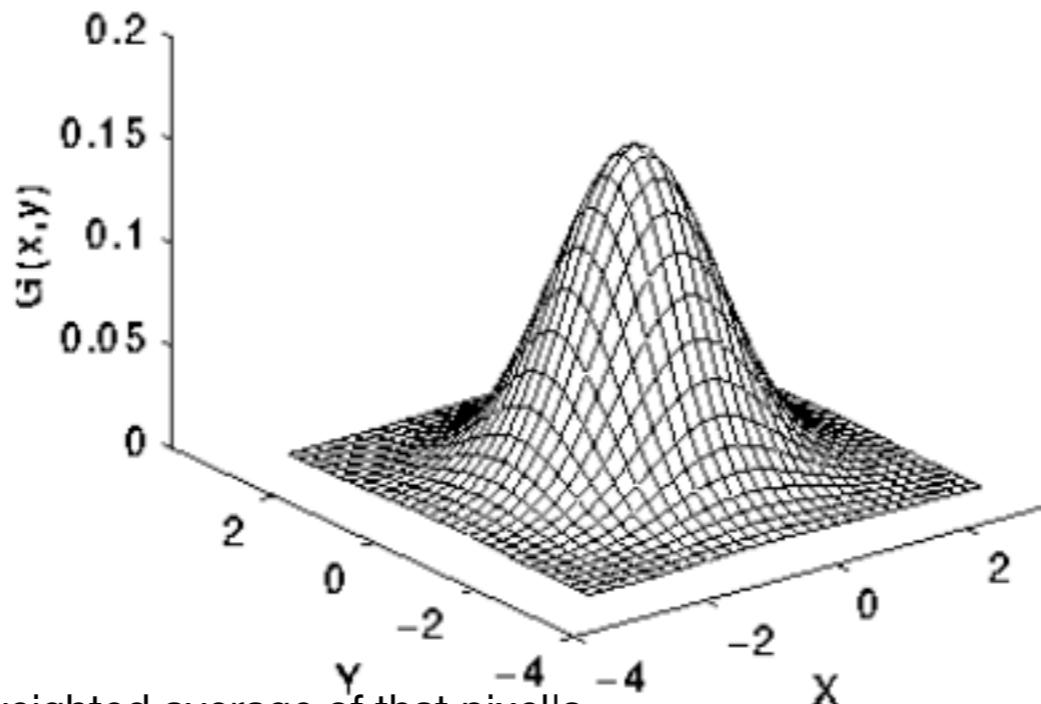
$$\theta(I, t) = \begin{cases} 1 & I > t \\ 0 & \text{otherwise} \end{cases}$$

the threshold value  
 $T(x,y)$   
is a weighted sum (cross-correlation with a Gaussian  
window) of the  
blockSize $\times$ blockSize  
neighborhood of  
( $x,y$ )  
minus C . The default sigma (standard deviation) is  
used for the specified blockSize

# Gaussian filter

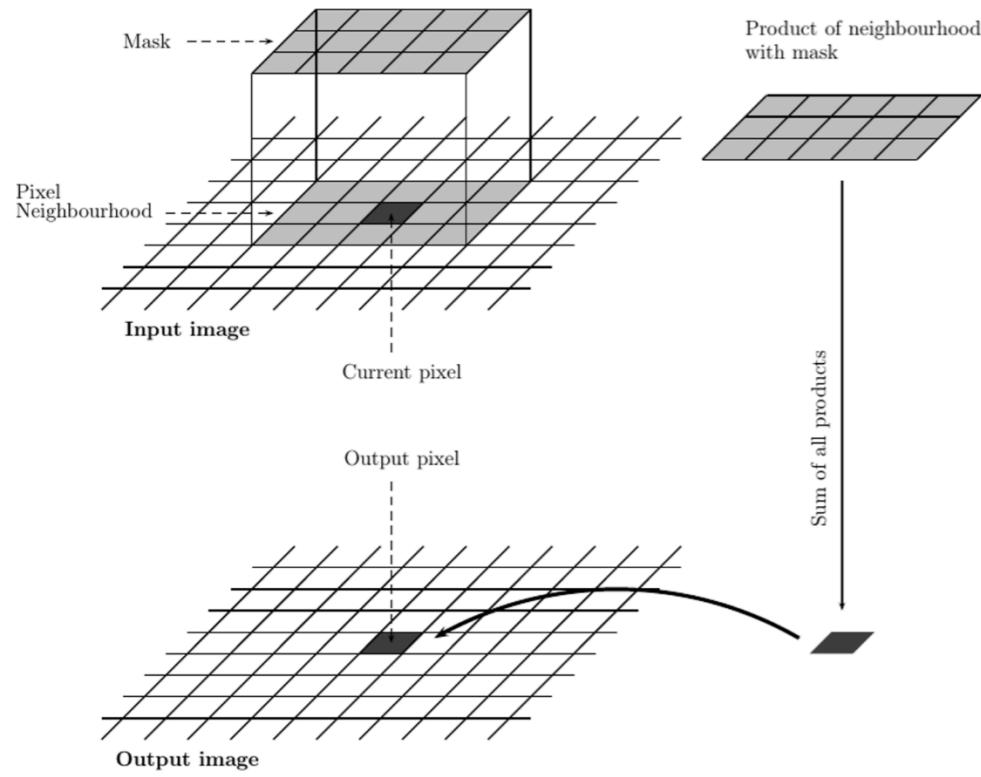
$$G(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$

$$\frac{1}{16} \begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix}$$



- Each pixel's new value is set to a weighted average of that pixel's neighborhood.
- The original pixel's value receives the heaviest weight and neighboring pixels receive smaller weights as their distance to the original pixel increases.
- This results in a blur that preserves boundaries and edges better than other, more uniform blurring filters
- 2 parameters: sigma and kernel radius





$$h[i, j] = \sum_{k=-a}^a \sum_{l=-b}^b f[k + a, l + b] \times I[i + k, j + l]$$

**f** = coefficient matrix, also known as kernel, filter, mask, ...

1/16

1	2	1
2	4	2
1	2	1

1/273

1	4	7	4	1
4	16	26	16	4
7	26	41	26	7
4	16	26	16	4
1	4	7	4	1

1/1003

0	0	1	2	1	0	0
0	3	13	22	13	3	0
1	13	59	97	59	13	1
2	22	97	159	97	22	2
1	13	59	97	59	13	1
0	3	13	22	13	3	0
0	0	1	2	1	0	0

# Thresholding and Binarization

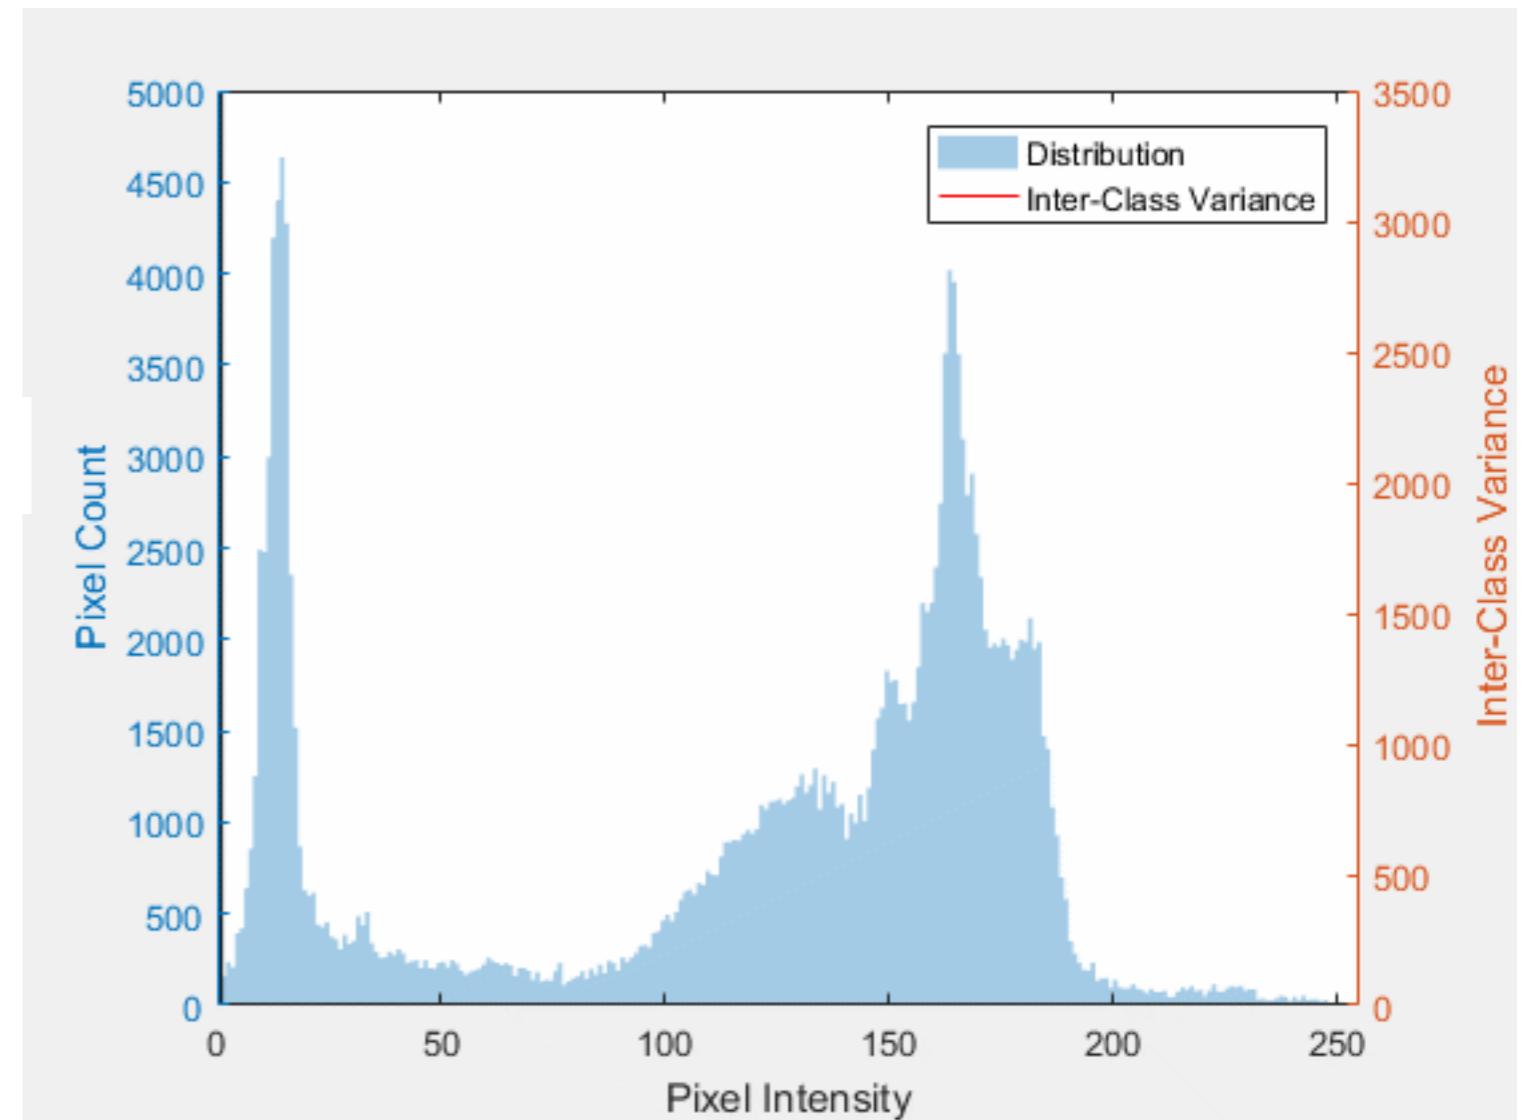
## Otsu Thresholding

$$\sigma_w^2(t) = \omega_0(t)\sigma_0^2(t) + \omega_1(t)\sigma_1^2(t)$$

$$\sigma_b^2(t) = \sigma^2 - \sigma_w^2(t) = \omega_1(t)\omega_2(t)[\mu_1(t) - \mu_2(t)]^2$$

Soglia migliore = quella che minimizza  $\sigma_w^2(t)$

Equivalent to maximizing  $\sigma_b^2(t)$



$$\omega_0(t) = \sum_{i=0}^{t-1} p(i)$$

Sigma = varianza dei punti  
sopra soglia e sotto soglia

$$\omega_1(t) = \sum_{i=t}^{L-1} p(i) \quad t=\text{soglia}$$

```

import numpy as np

def compute_otsu_criteria(im, th):
    # create the thresholded image
    thresholded_im = np.zeros(im.shape)
    thresholded_im[im >= th] = 1

    # compute weights
    nb_pixels = im.size
    nb_pixels1 = np.count_nonzero(thresholded_im)
    weight1 = nb_pixels1 / nb_pixels
    weight0 = 1 - weight1

    # if one the classes is empty, eg all pixels are below or above the threshold, that threshold will not be considered
    # in the search for the best threshold
    if weight1 == 0 or weight0 == 0:
        return np.inf

    # find all pixels belonging to each class
    val_pixels1 = im[thresholded_im == 1]
    val_pixels0 = im[thresholded_im == 0]

    # compute variance of these classes
    var0 = np.var(val_pixels0) if len(val_pixels0) > 0 else 0
    var1 = np.var(val_pixels1) if len(val_pixels1) > 0 else 0

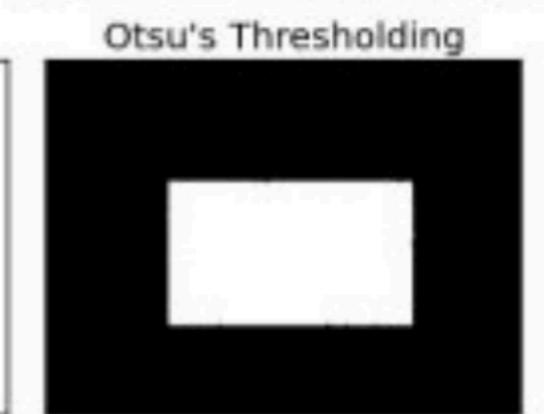
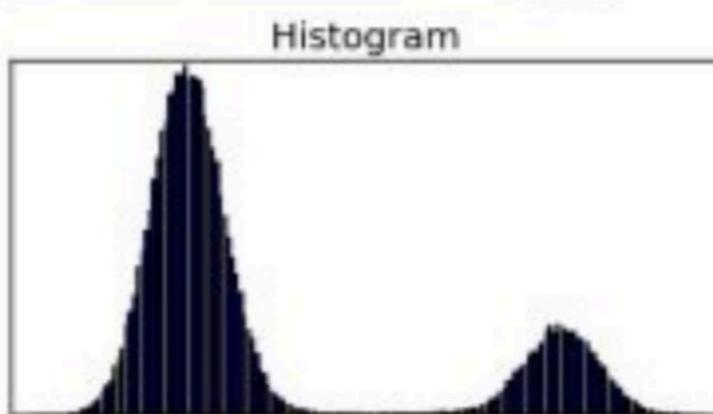
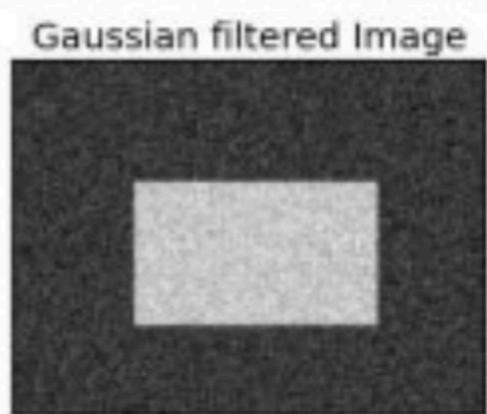
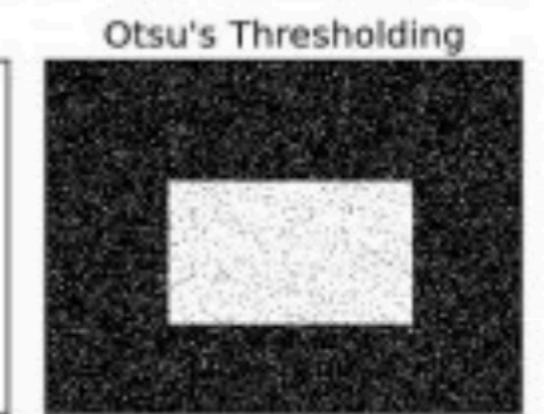
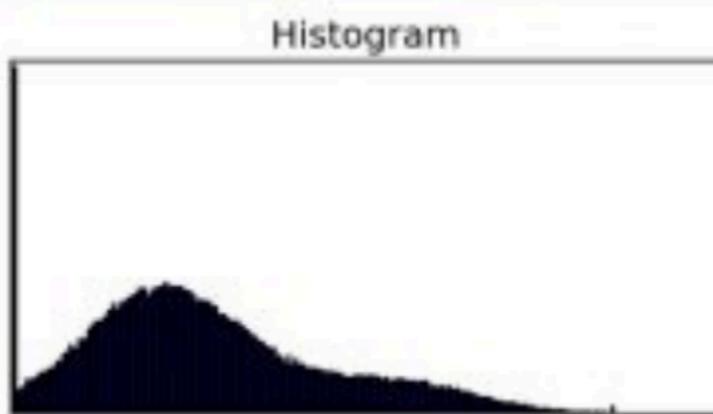
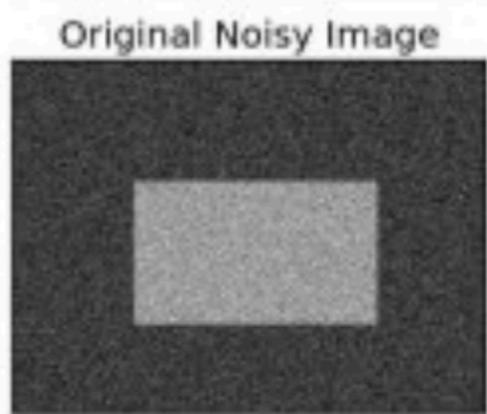
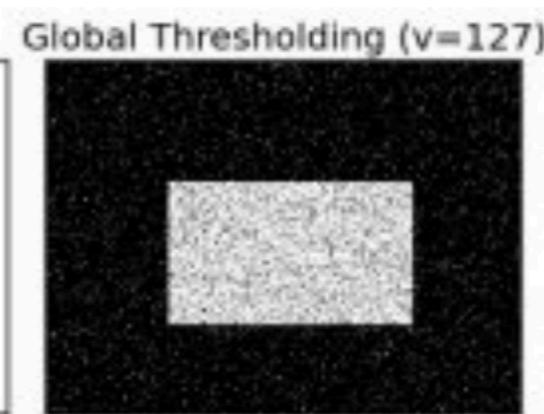
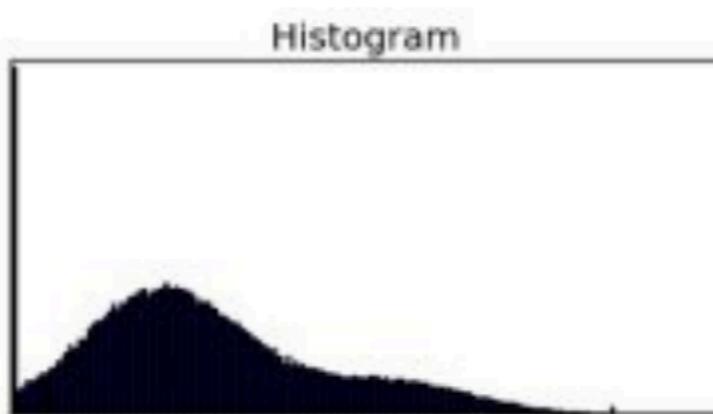
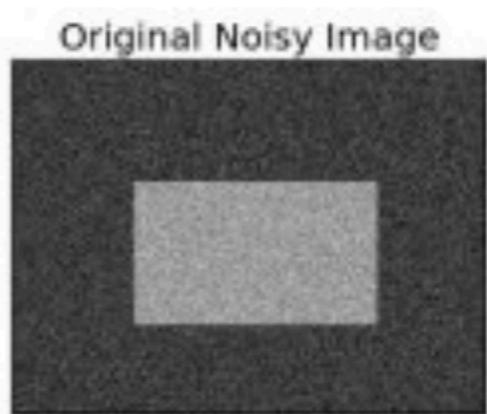
    return weight0 * var0 + weight1 * var1

im = # load your image as a numpy array.
# For testing purposes, one can use for example im = np.random.randint(0,255, size = (50,50))

# testing all thresholds from 0 to the maximum of the image
threshold_range = range(np.max(im)+1)
criterias = [compute_otsu_criteria(im, th) for th in threshold_range]

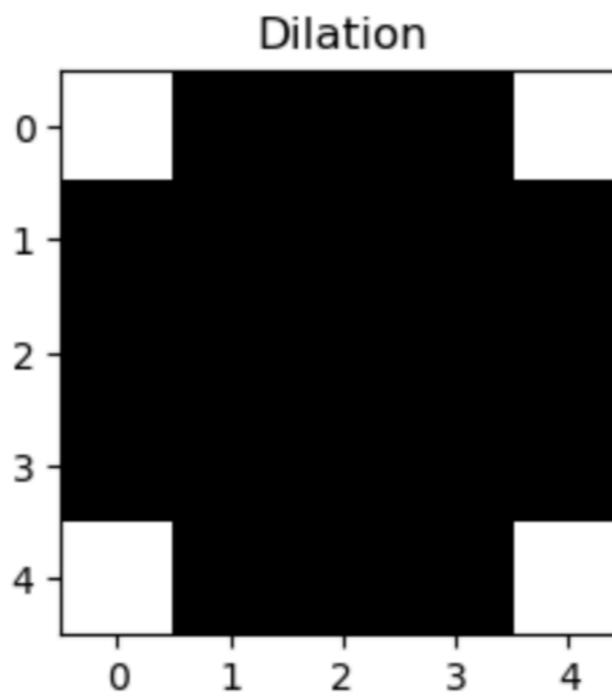
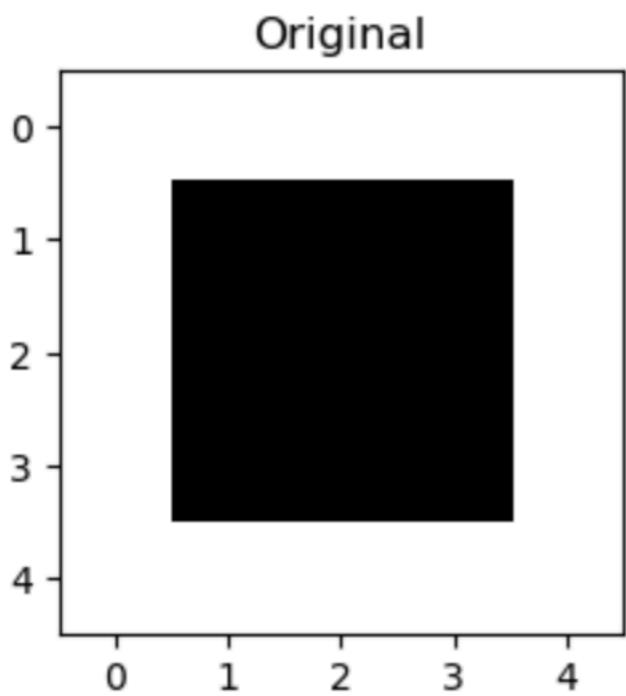
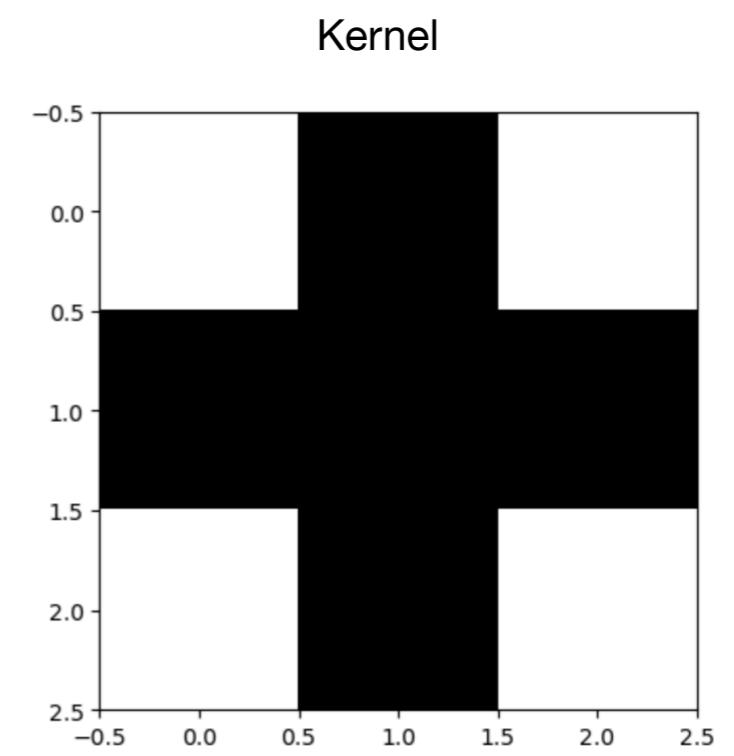
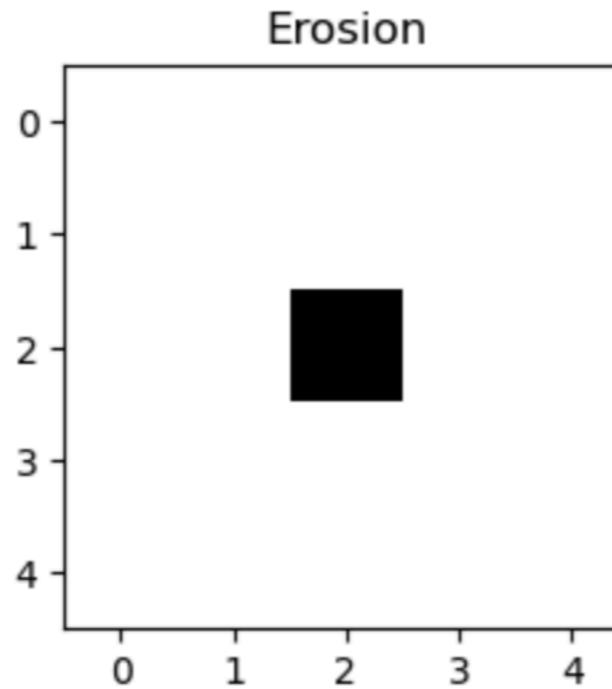
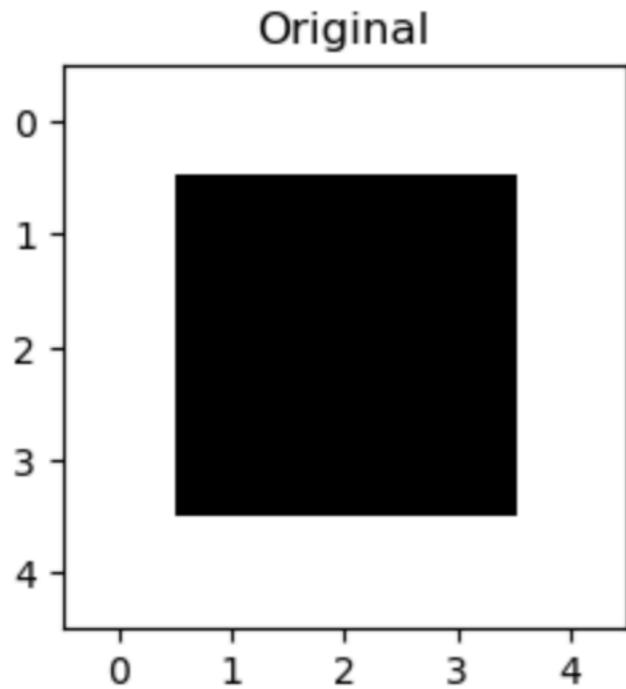
# best threshold is the one minimizing the Otsu criteria
best_threshold = threshold_range[np.argmin(criterias)]

```



**Demo dei vari thresholding su jupyter  
th\_demo.ipynb**

# Erosion and dilation



Erosion=minimum  
Dilation = maximum

# **Opening and closing**

**Opening = Erosion+ Dilation**



**Closing = Dilation +Erosion**



# First segmentation

Thresholding + closing + binarization + contour

Contour = curve joining all the continuous points with same intensity



# Edge detection

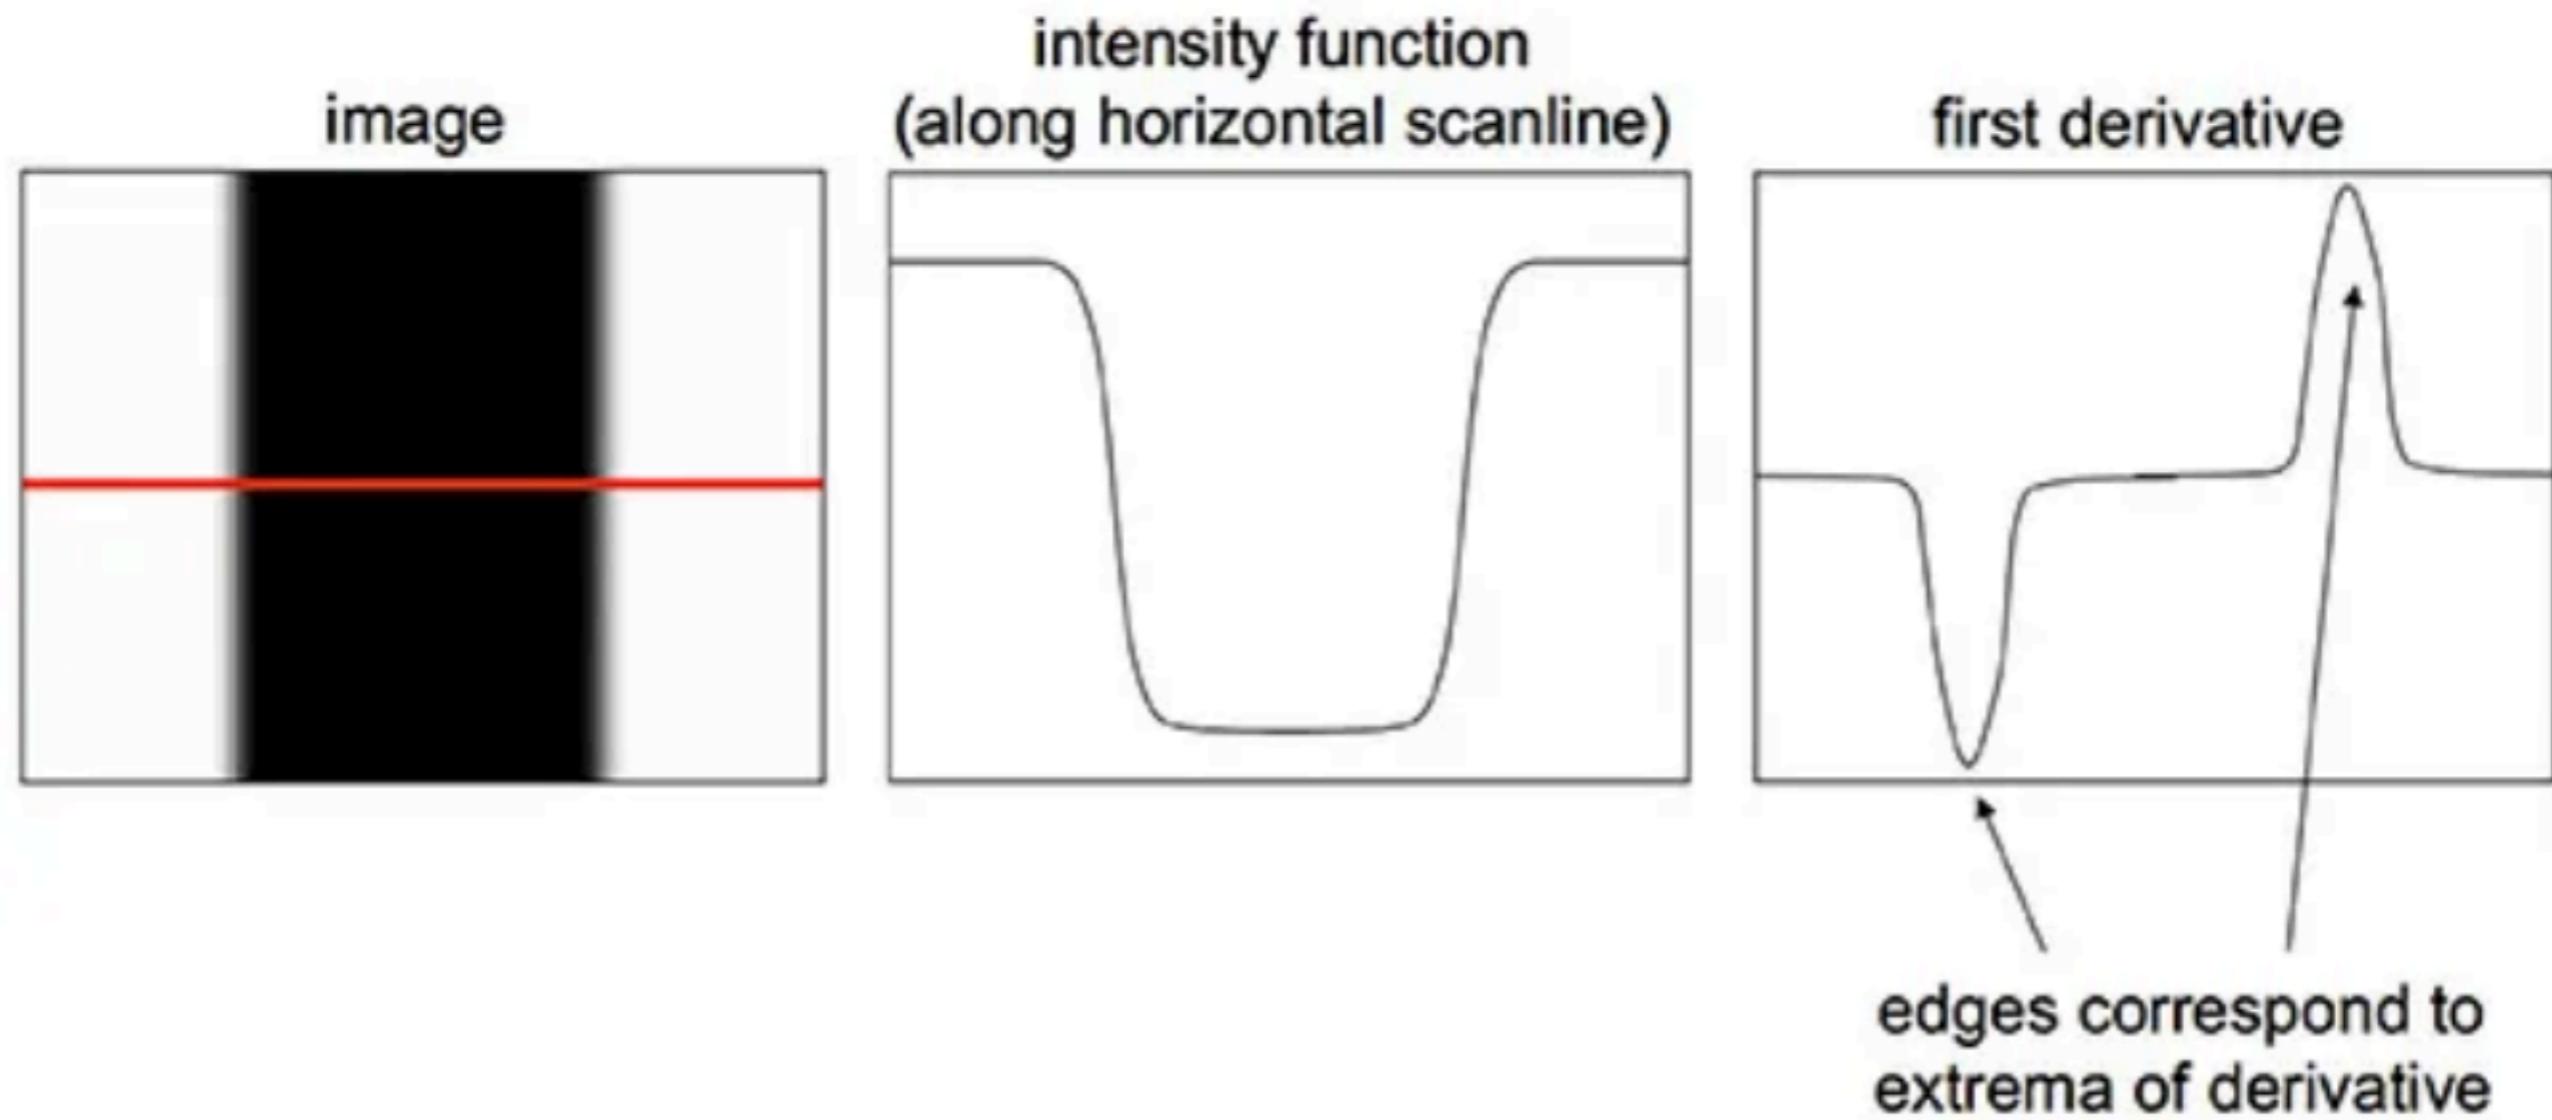
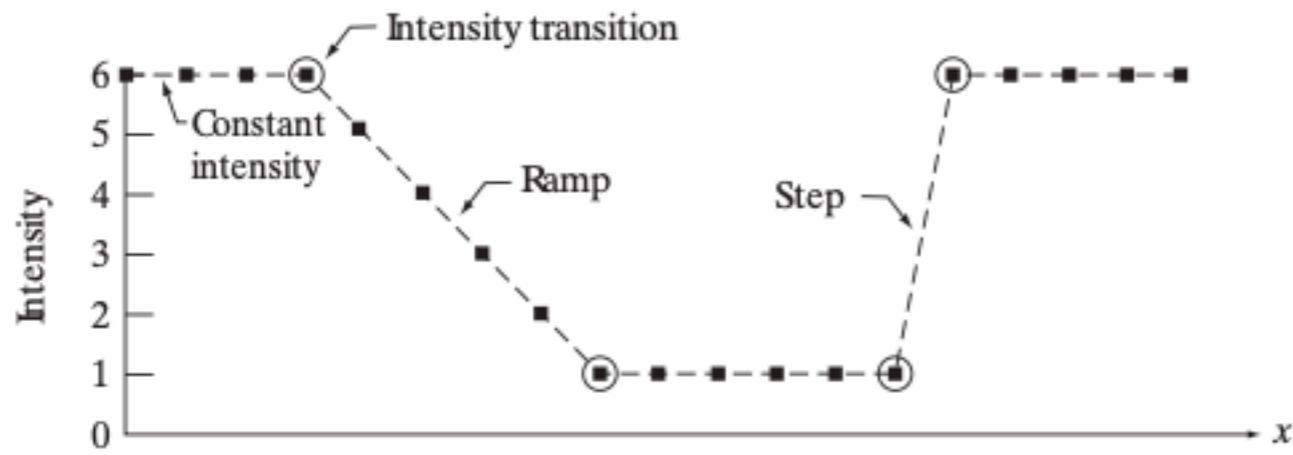


Image Credit: <http://stanford.edu/>



Scan line

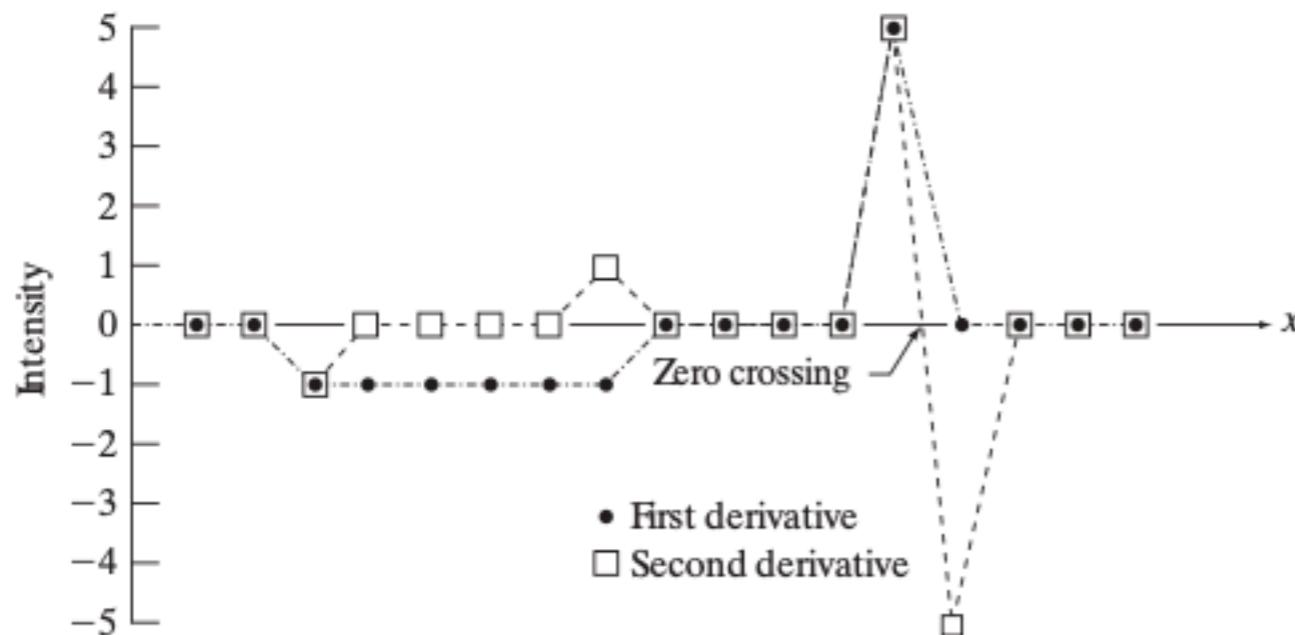
6	6	6	6	5	4	3	2	1	1	1	1	1	1	6	6	6	6	6
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

1st derivative

0	0	-1	-1	-1	-1	0	0	0	0	0	0	5	0	0	0	0	0
---	---	----	----	----	----	---	---	---	---	---	---	---	---	---	---	---	---

2nd derivative

0	0	-1	0	0	0	0	1	0	0	0	0	5	-5	0	0	0	0
---	---	----	---	---	---	---	---	---	---	---	---	---	----	---	---	---	---



$$\frac{\partial f}{\partial x} = f(x+1) - f(x)$$

$$\frac{\partial^2 f}{\partial x^2} = f(x+1) + f(x-1) - 2f(x)$$

# Sobel filter

**Orizzontale**

1	0	-1
2	0	-2
1	0	-1

=

1
2
1

\*

1	0	-1
---	---	----

**Verticale**

1	2	1
0	0	0
-1	-2	-1

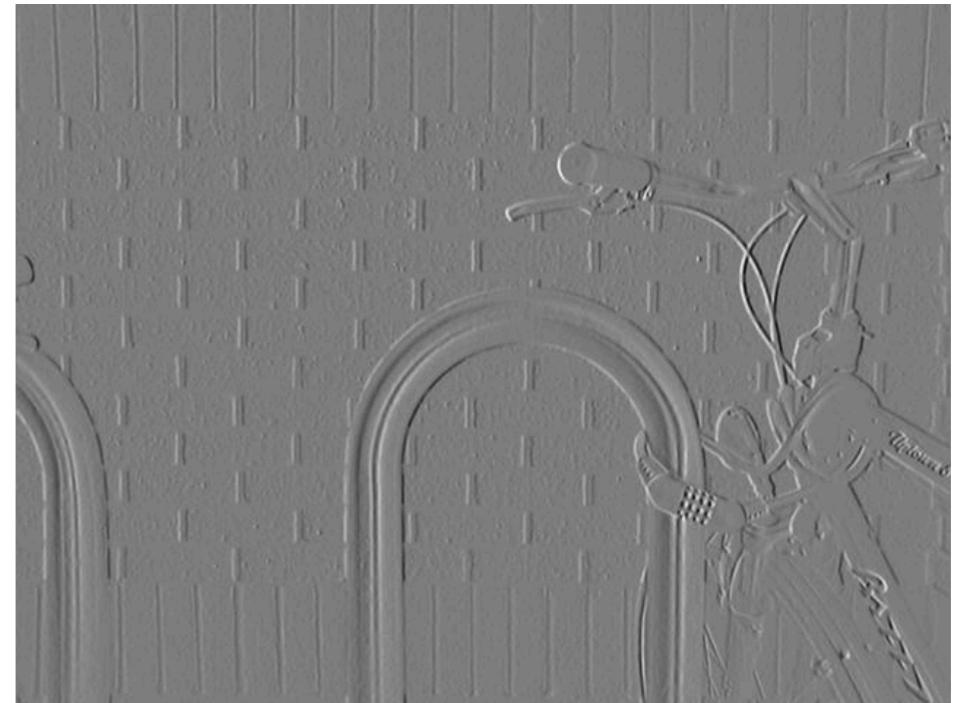
=

1
0
-1

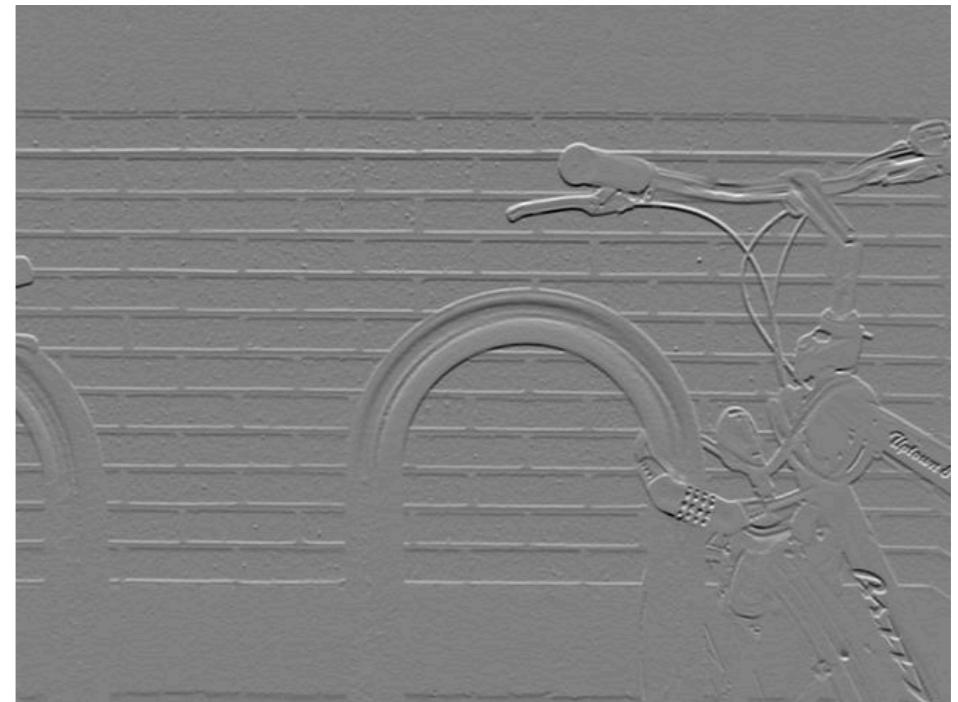
\*

1	2	1
---	---	---

# Orizzontale



# Verticale



# Gradient-based filtering

1. Select your favorite derivative filters.

$$\mathbf{S}_x = \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 2 & 0 & -2 \\ \hline 1 & 0 & -1 \\ \hline \end{array}$$

$$\mathbf{S}_y = \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array}$$

2. Convolve with the image

$$\frac{\partial f}{\partial x} = \mathbf{S}_x * f$$

$$\frac{\partial f}{\partial y} = \mathbf{S}_y * f$$

3. Calculate direction and magnitude of the gradient.

$$\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

gradient

$$\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

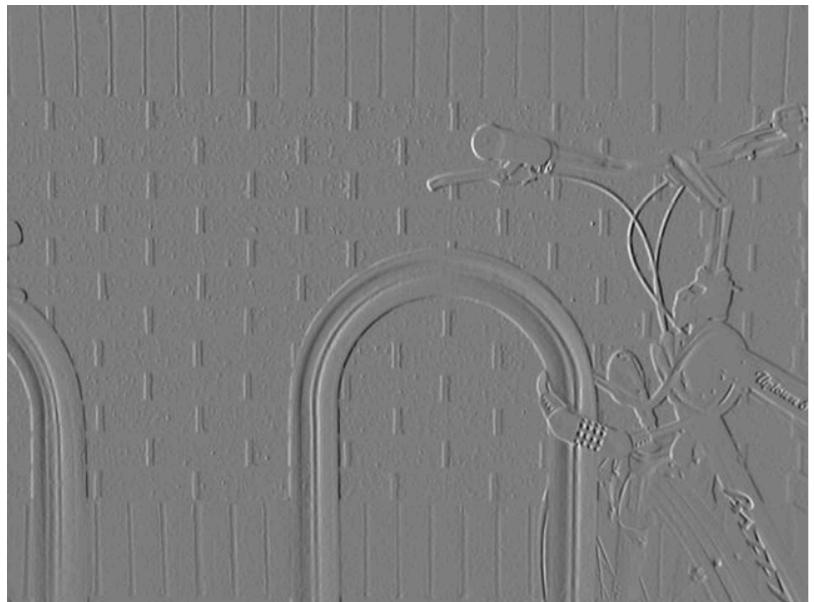
direction

$$||\nabla f|| = \sqrt{\left( \frac{\partial f}{\partial x} \right)^2 + \left( \frac{\partial f}{\partial y} \right)^2}$$

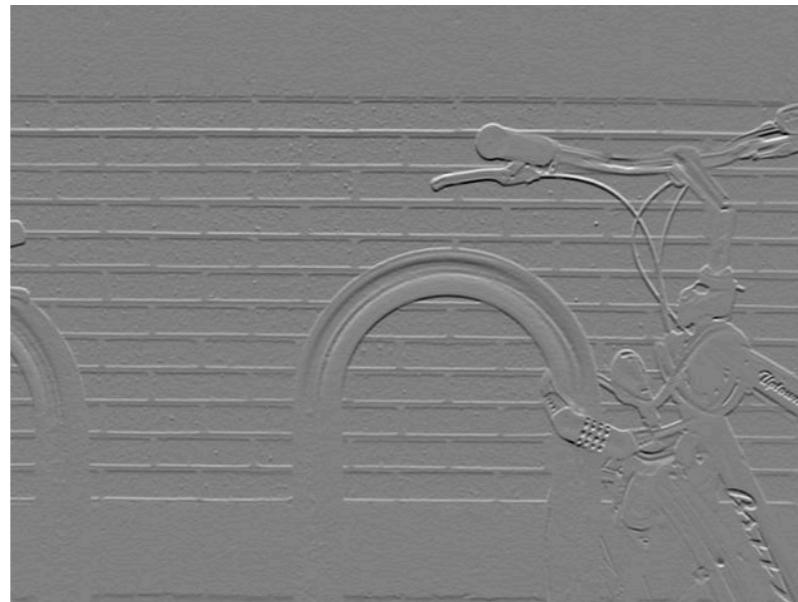
amplitude



**Orizzontale**



**Verticale**



**Gradient**



# **Demo sobel**

## **Edge detection - Canny**

- Edges with gradient could be good, but typically some edges thick than others
- Ideally all edges should be thin and with the same intensity (255)

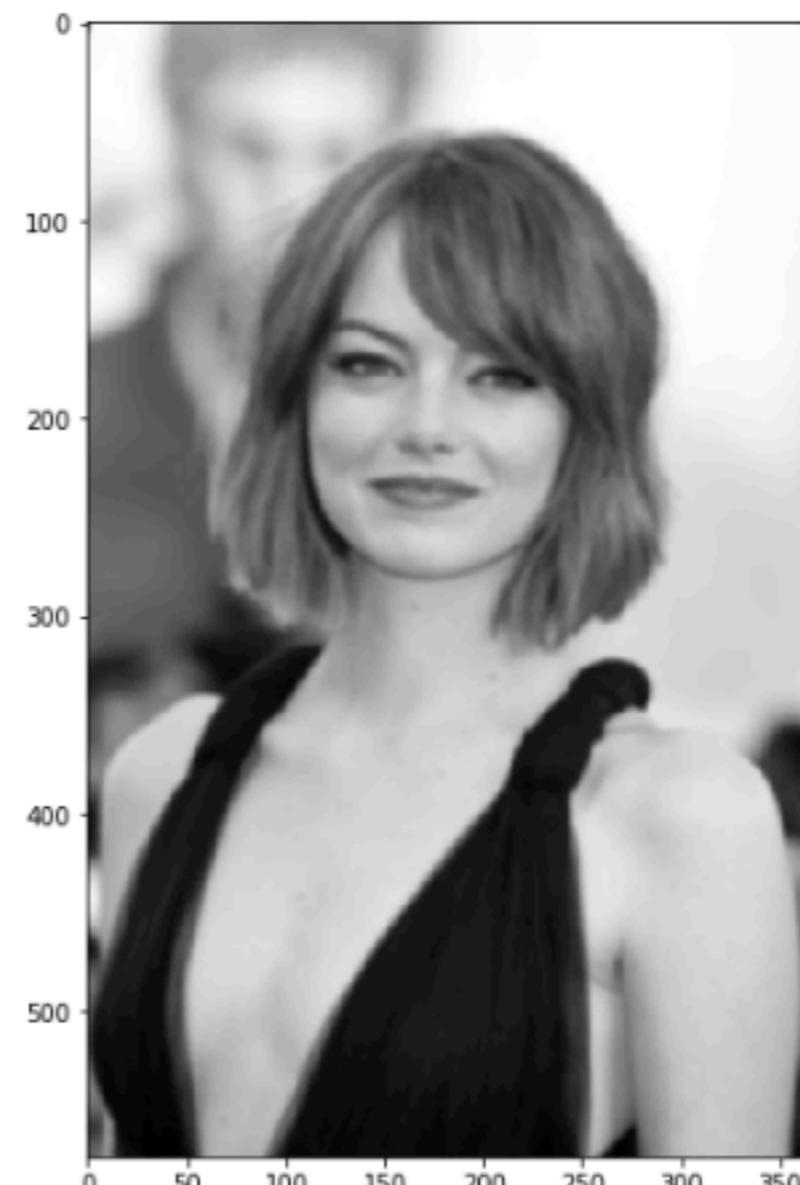
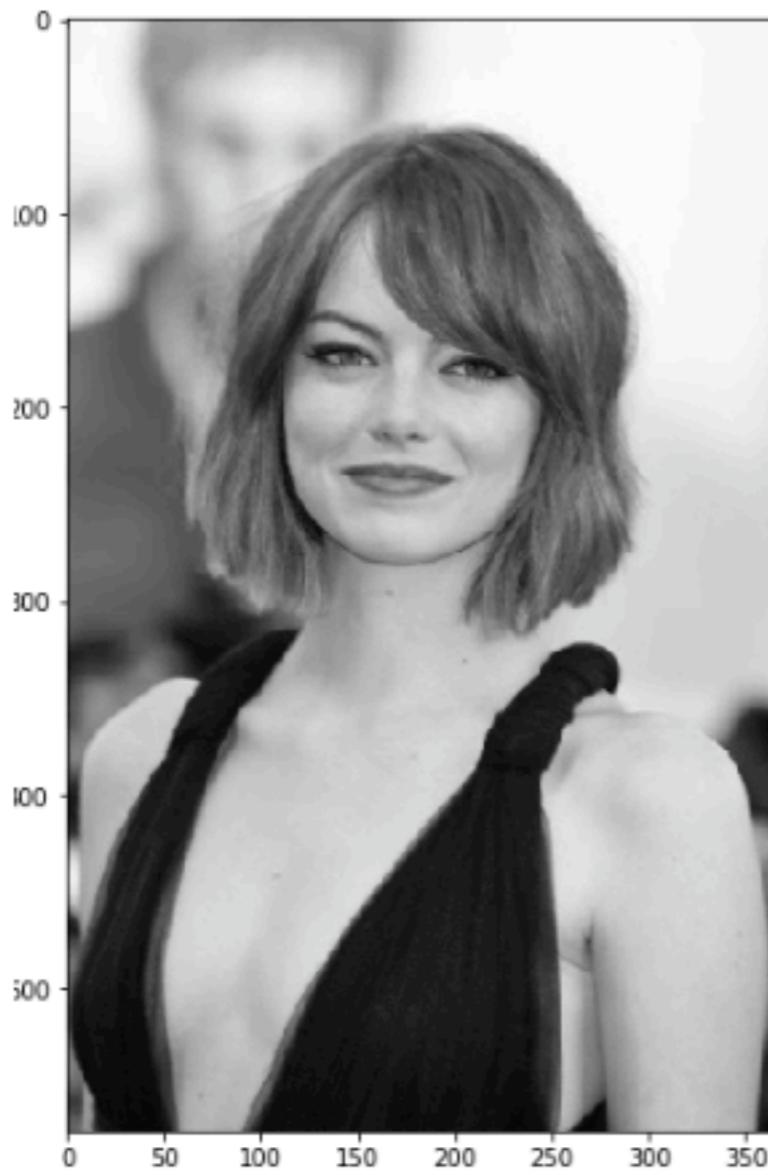
# **Edge detection - Canny**

## **5 Steps:**

1. Noise reduction
2. Gradient calculation
3. Non-maximum suppression
4. Double threshold
5. Edge Tracking by Hysteresis

# 1. Noise reduction

## Gaussian blur



Original image (left) — Blurred image with a Gaussian filter (sigma=1.4 and kernel size of 5×5)

## 2. Gradient calculation

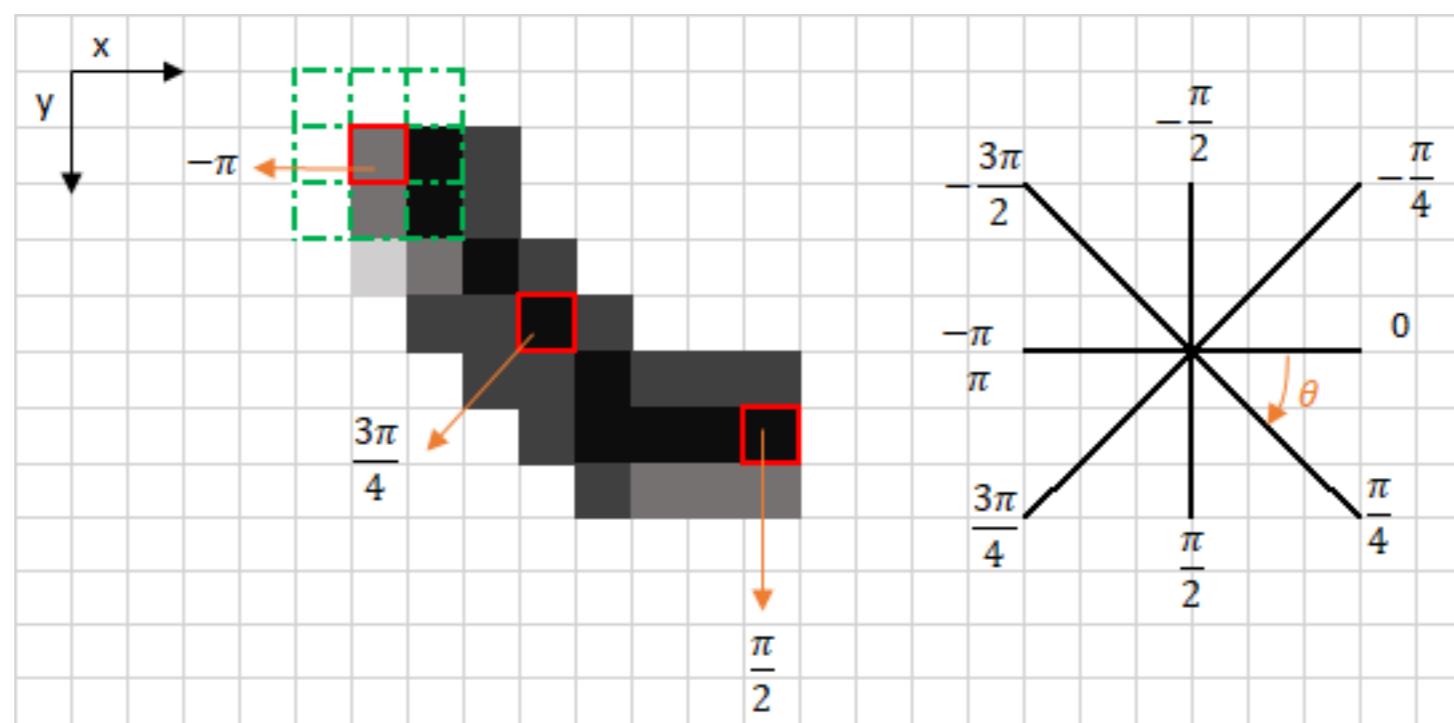


Blurred image (left) — Gradient intensity (right)

### 3. Non-maximum suppression

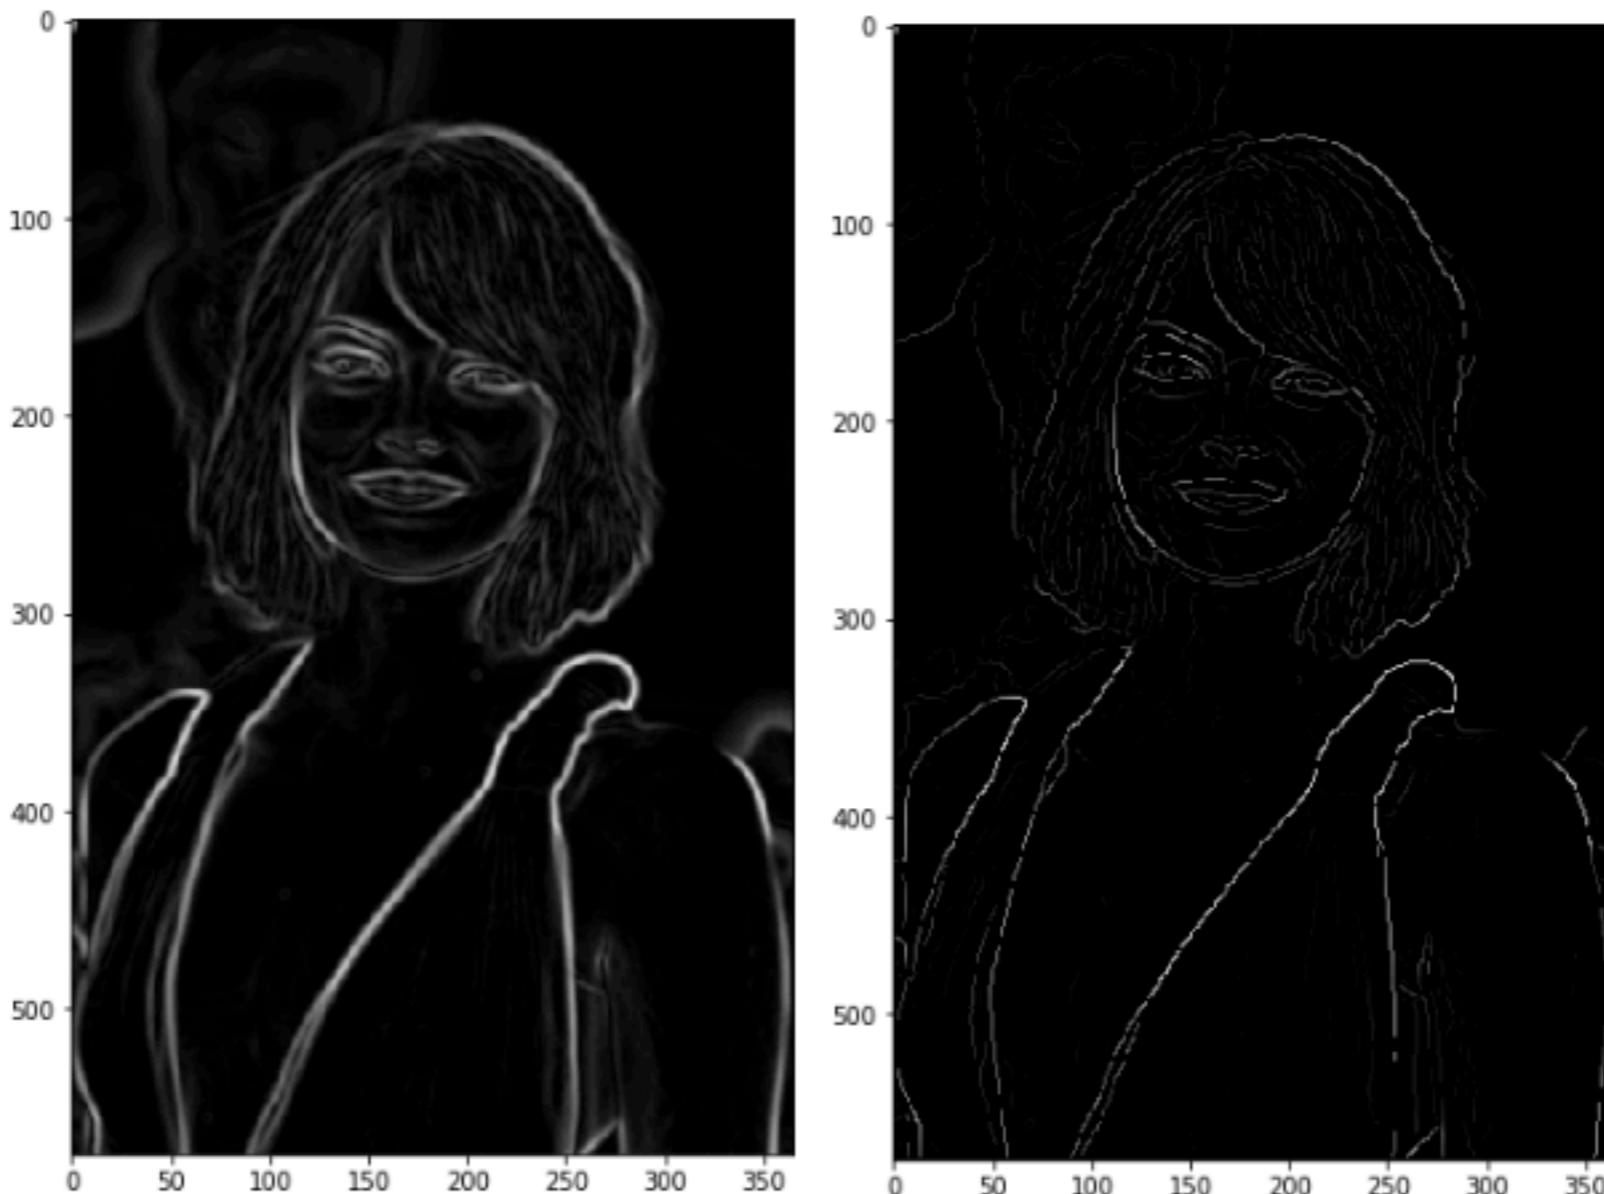
$$\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

direction



Pixels with a first neighbor in theta direction with bigger values are set to zero

### 3. Non-maximum suppression



Result of the non-max suppression.

Now thinner but with different intensities

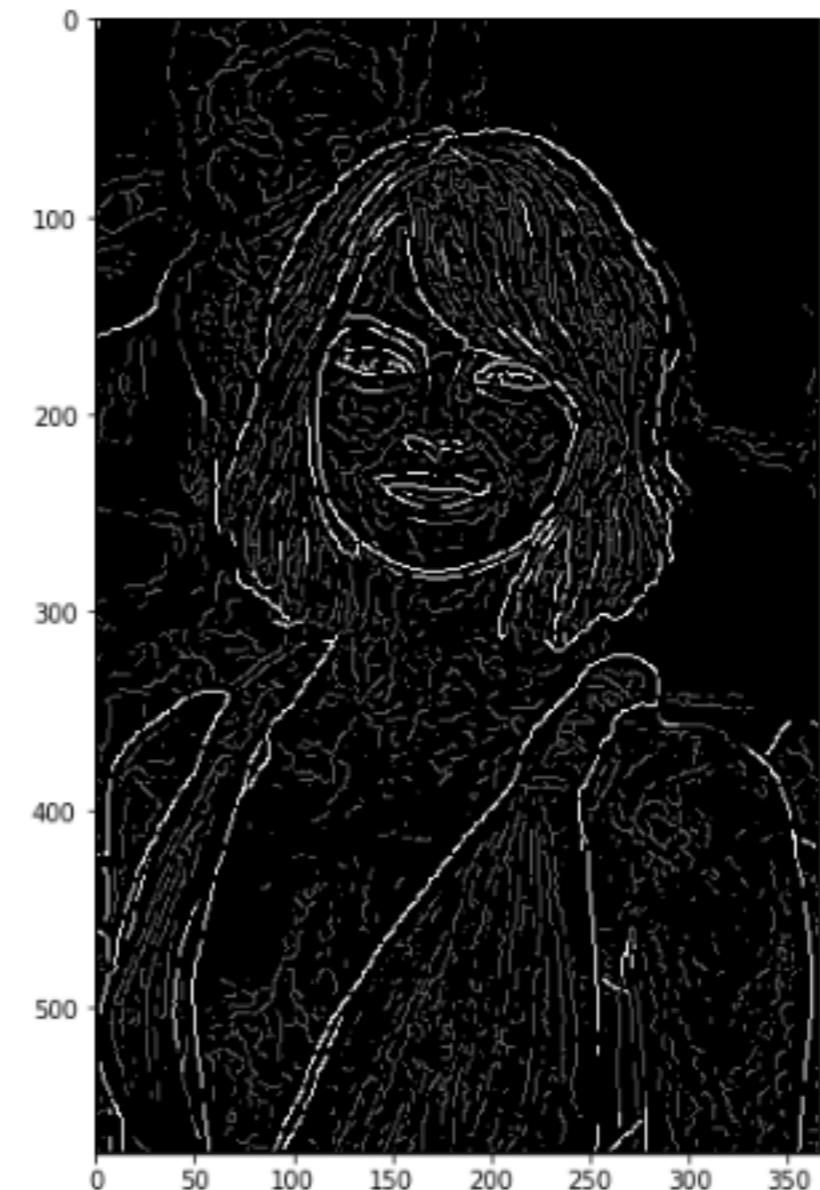
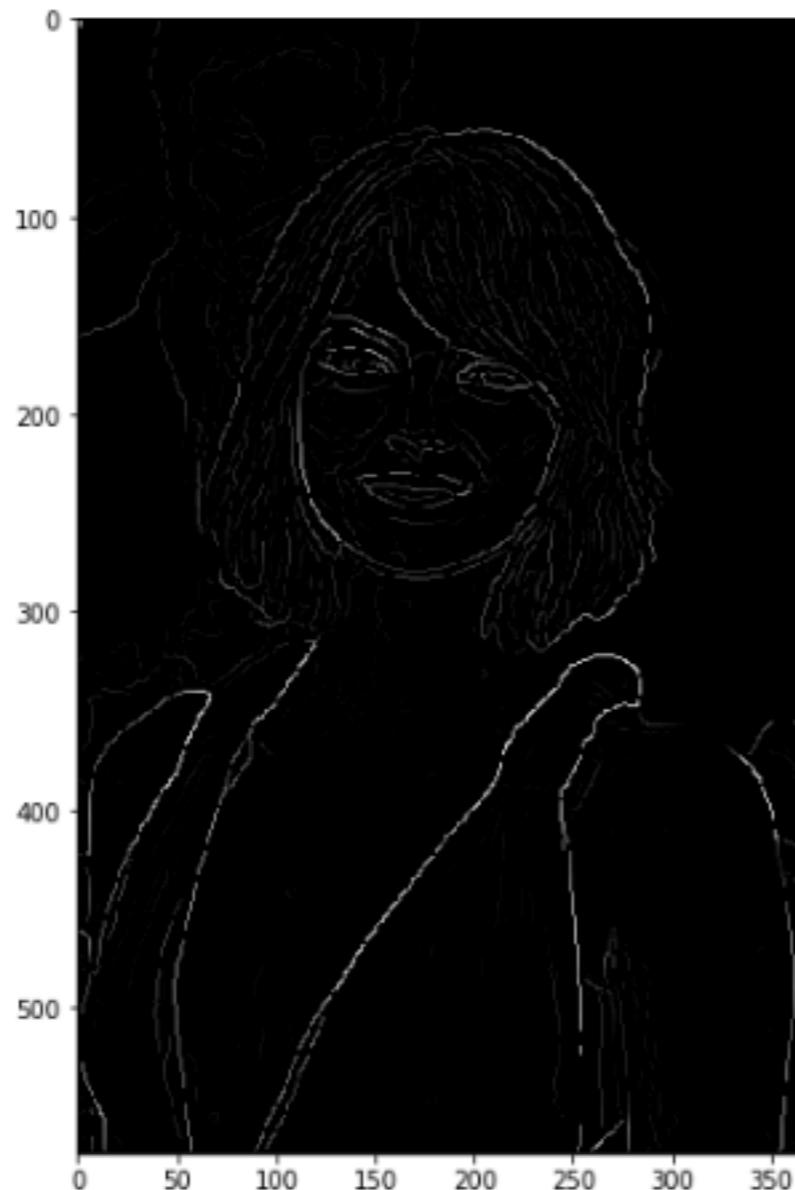
## 4. Double threshold

Two thresholds: high and low (es. 70,10)

**Strong** pixel:  $\text{img}(x,y) \geq \text{high} \rightarrow \text{img}(x,y) = 255$

**Weak** pixel:  $\text{img}(x,y) < \text{high}$  **and**  $\text{img}(x,y) \geq \text{low} \rightarrow \text{img}(x,y) = 25$

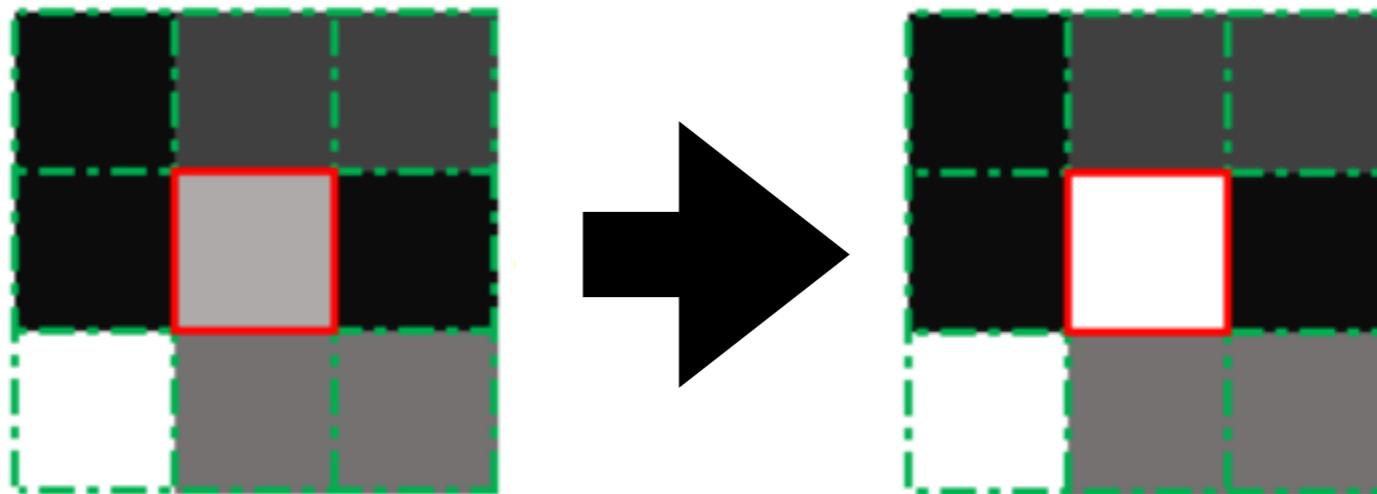
**Zero** pixel:  $\text{img}(x,y) < \text{low} \rightarrow \text{img}(x,y) = 0$



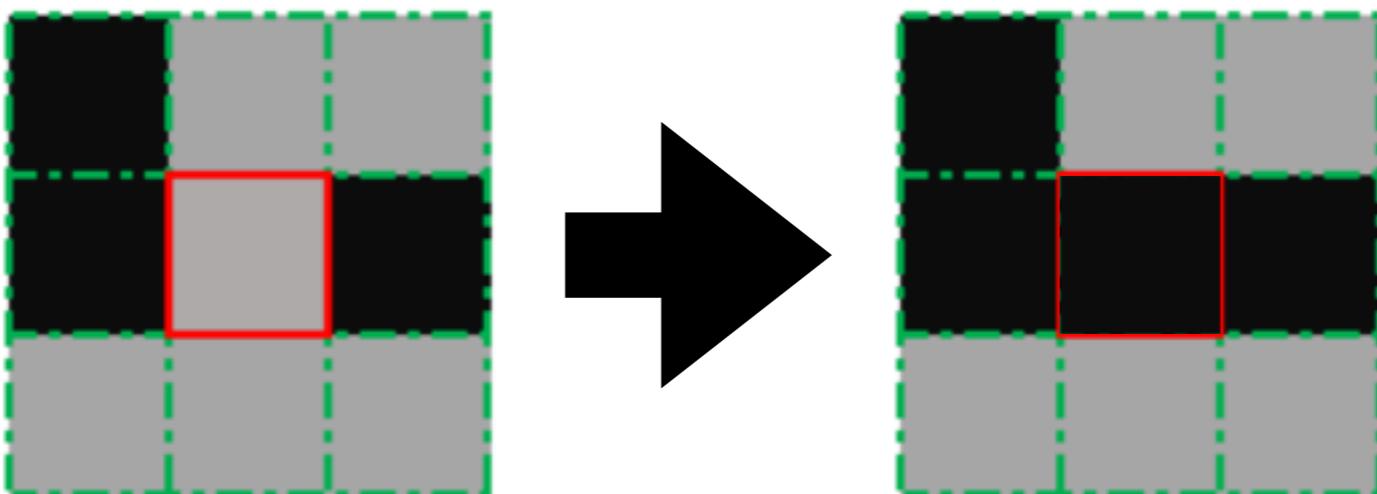
Non-Max Suppression image (left) — Threshold result (right): weak pixels in gray and strong ones in white.

## 5. Edge Tracking by Hysteresis

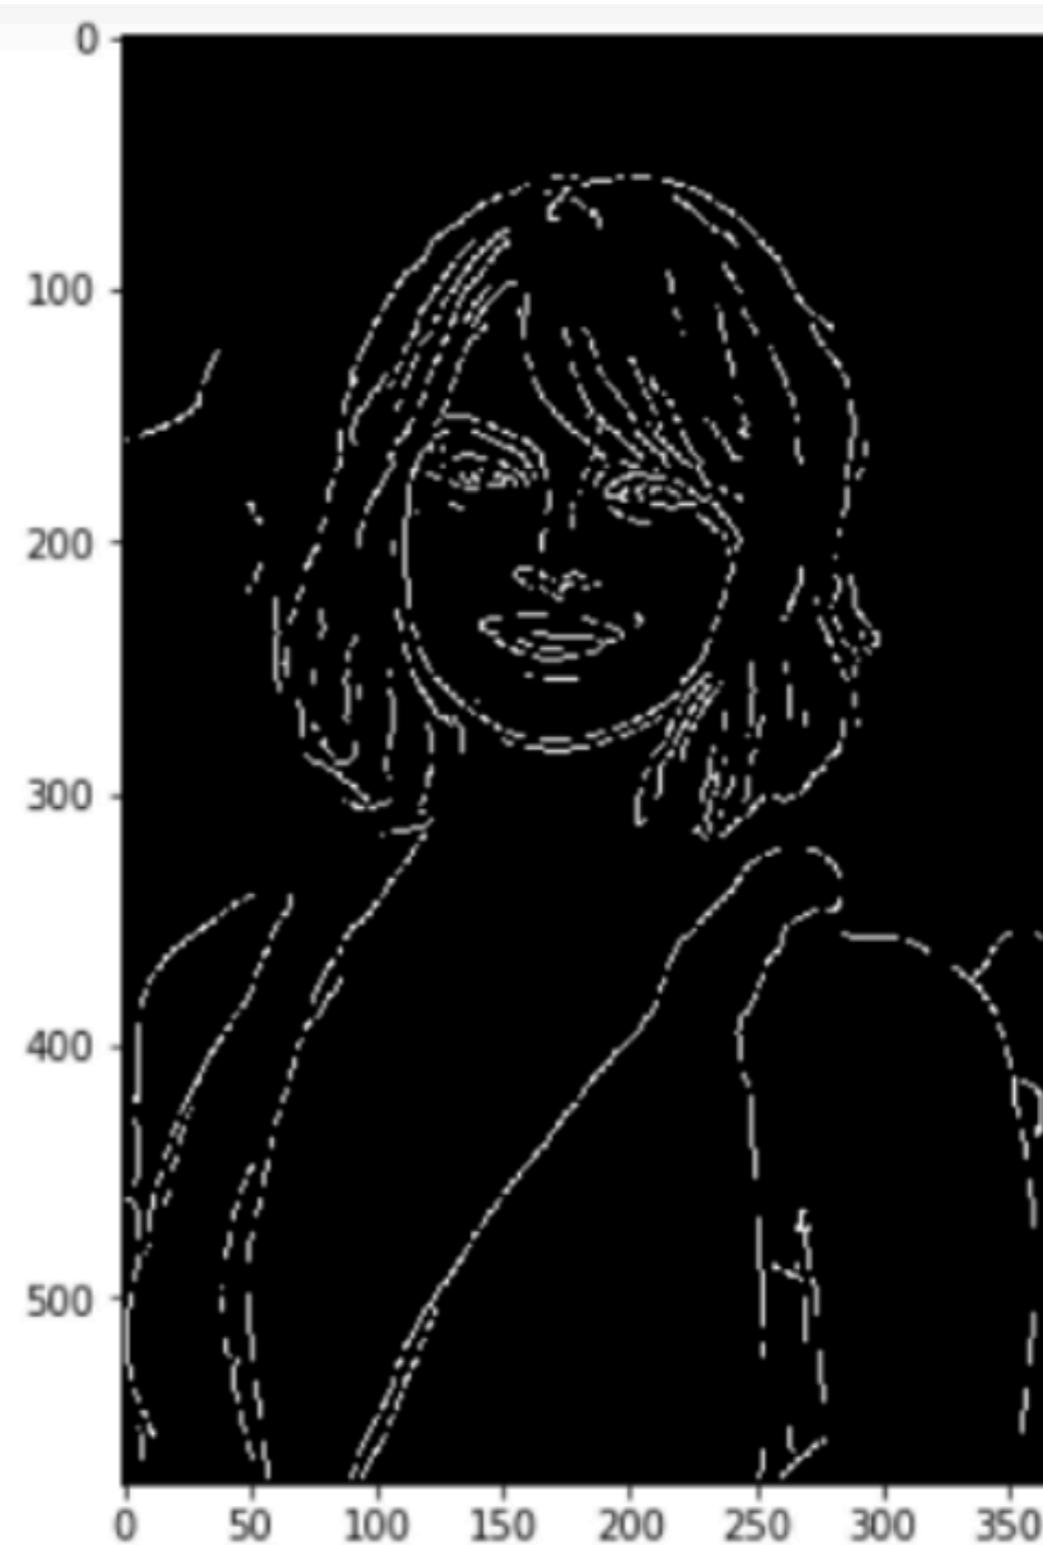
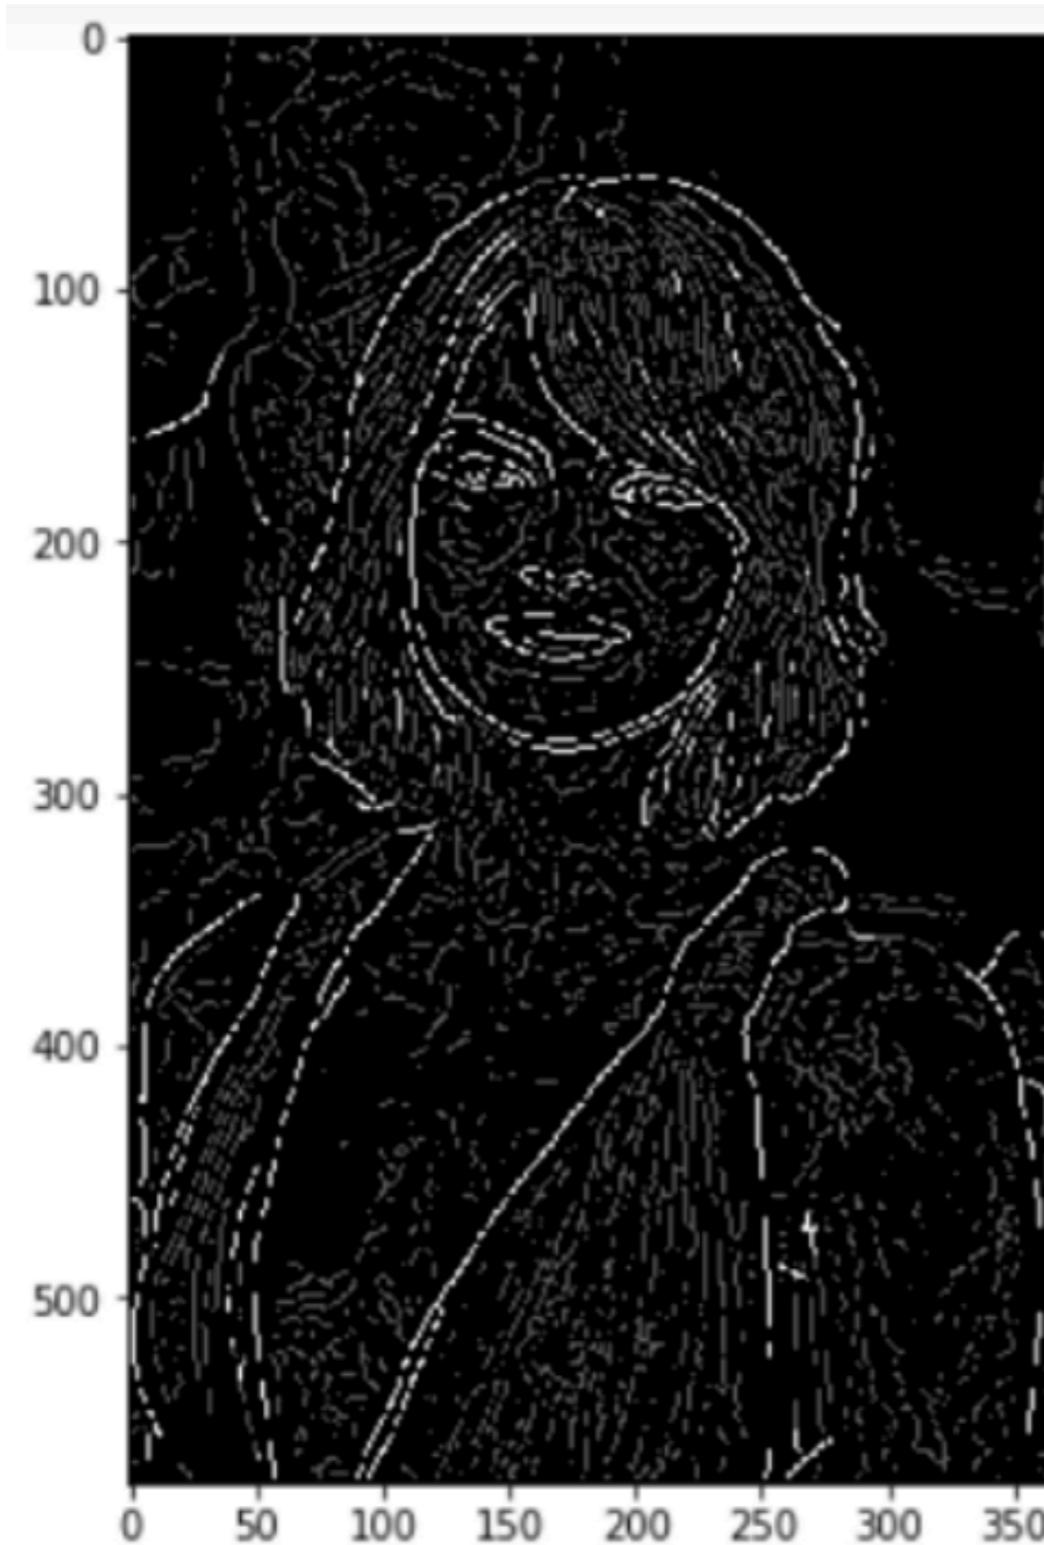
Strong pixel around



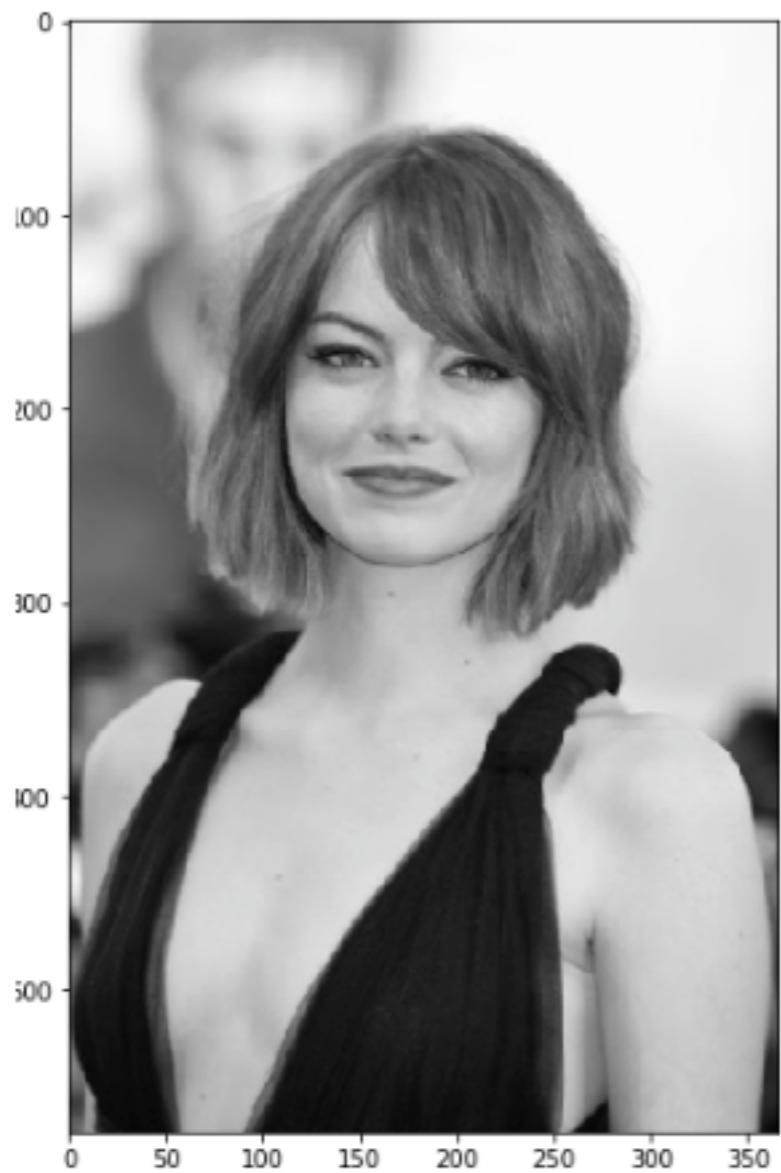
NO Strong pixel around



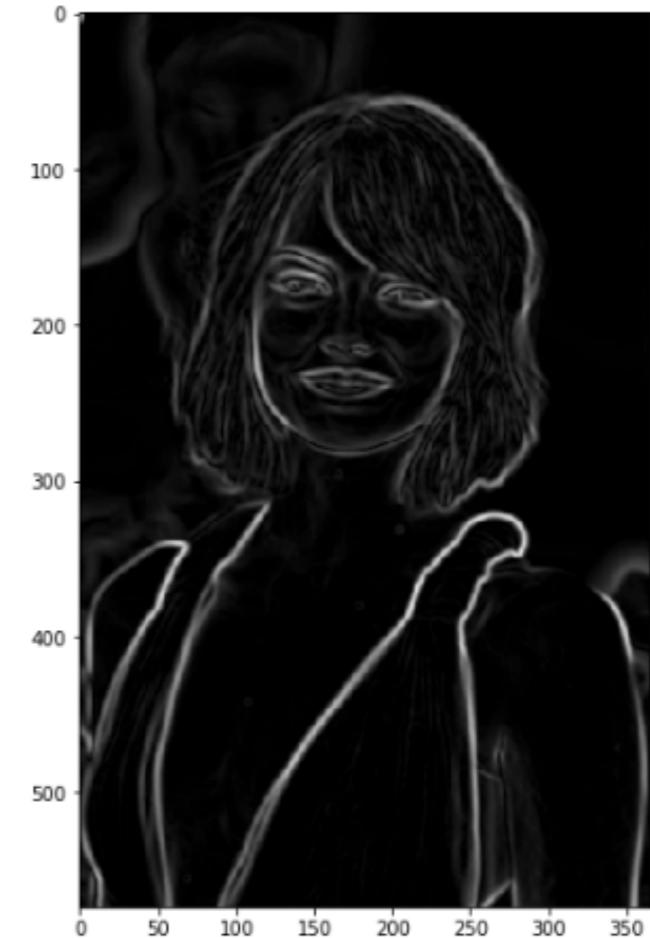
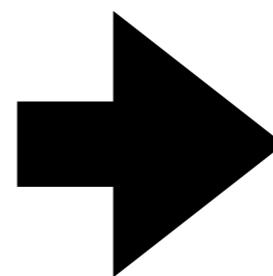
## 5. Edge Tracking by Hysteresis



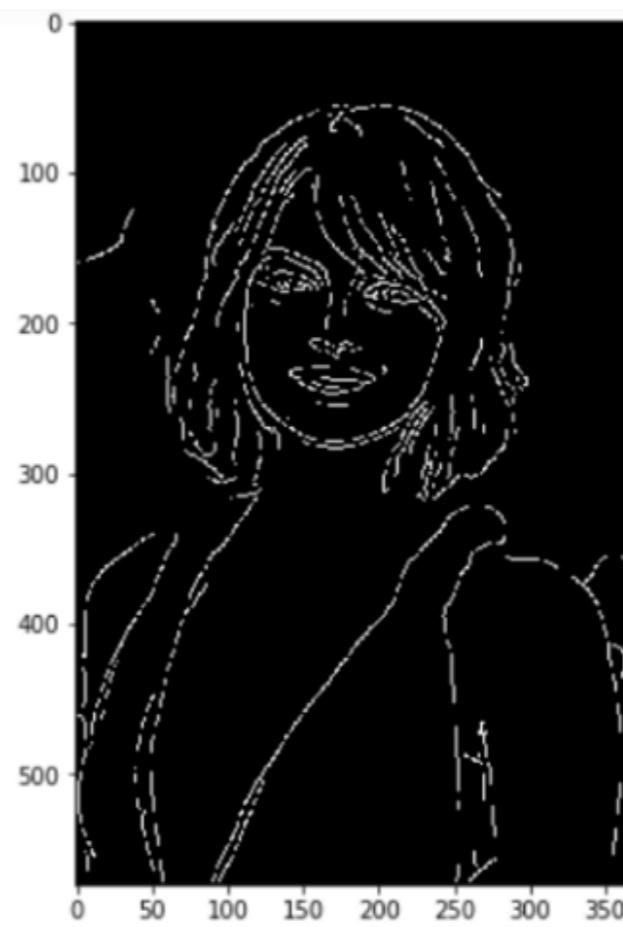
Results of hysteresis process



**Gradient (Sobel)**



**Canny**



# **Demo canny**

## Second segmentation

Edge detection + contour



# Watershed segmentation

Threshold segmentation



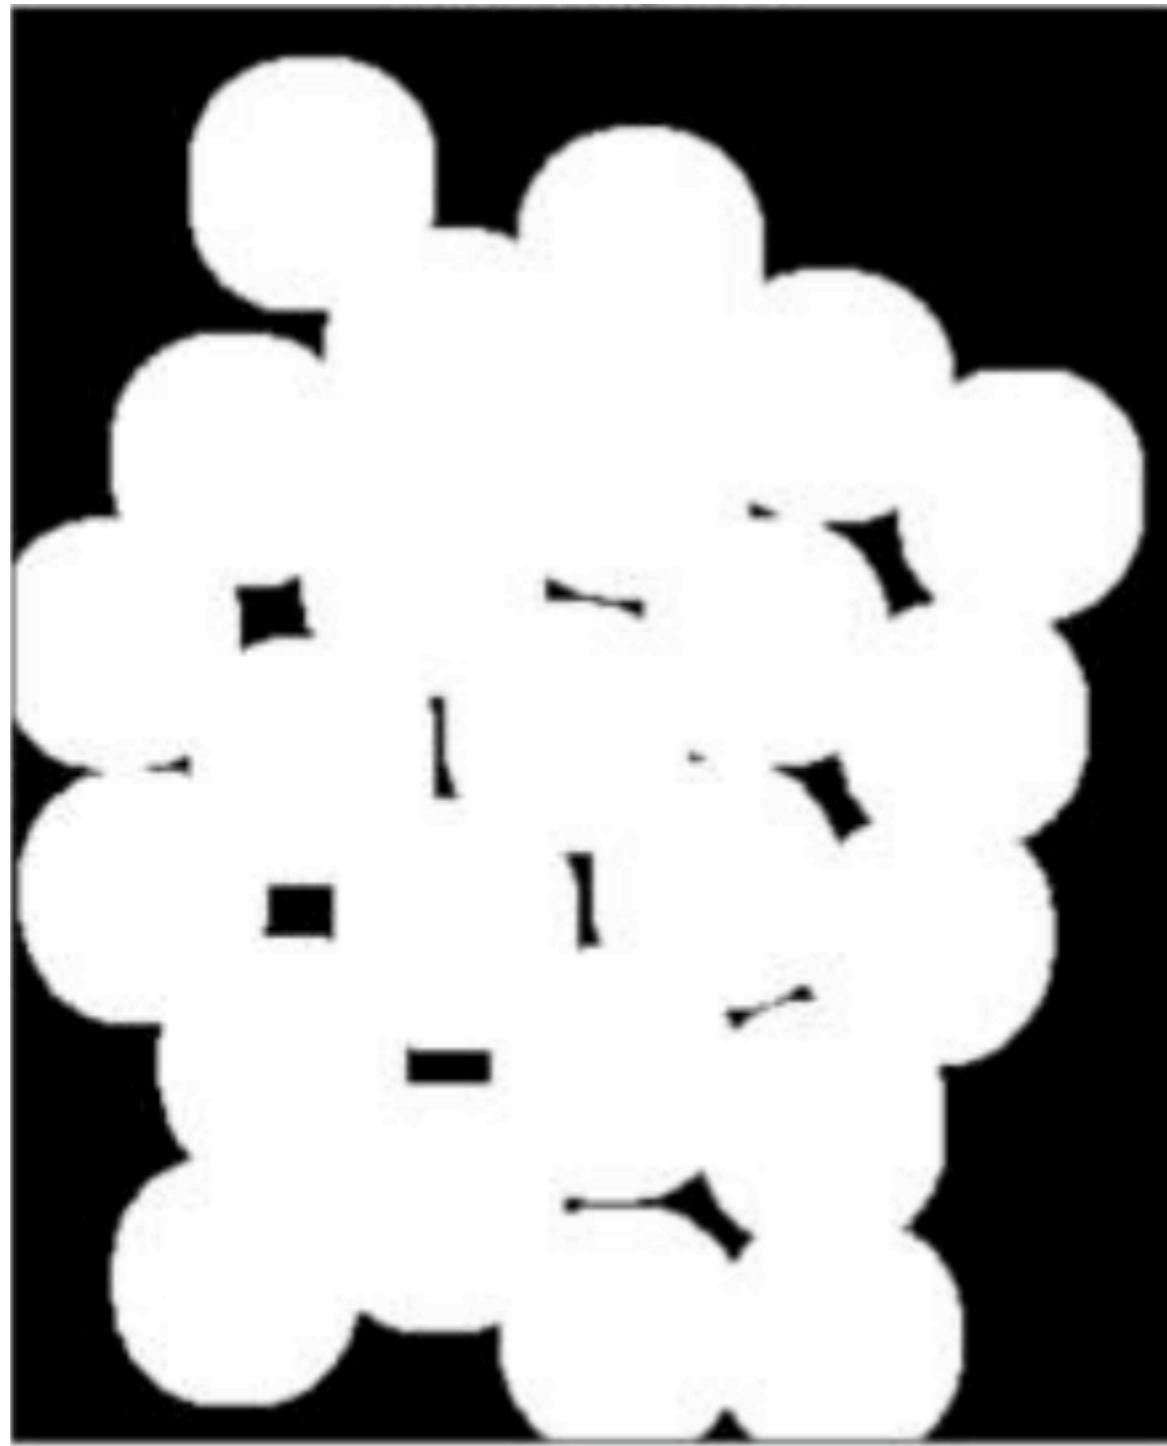
Edge segmentation



**Contours are touching each other, how to separate the coins contours ?**

# Watershed segmentation

Otsu thresholding + opening + dilate

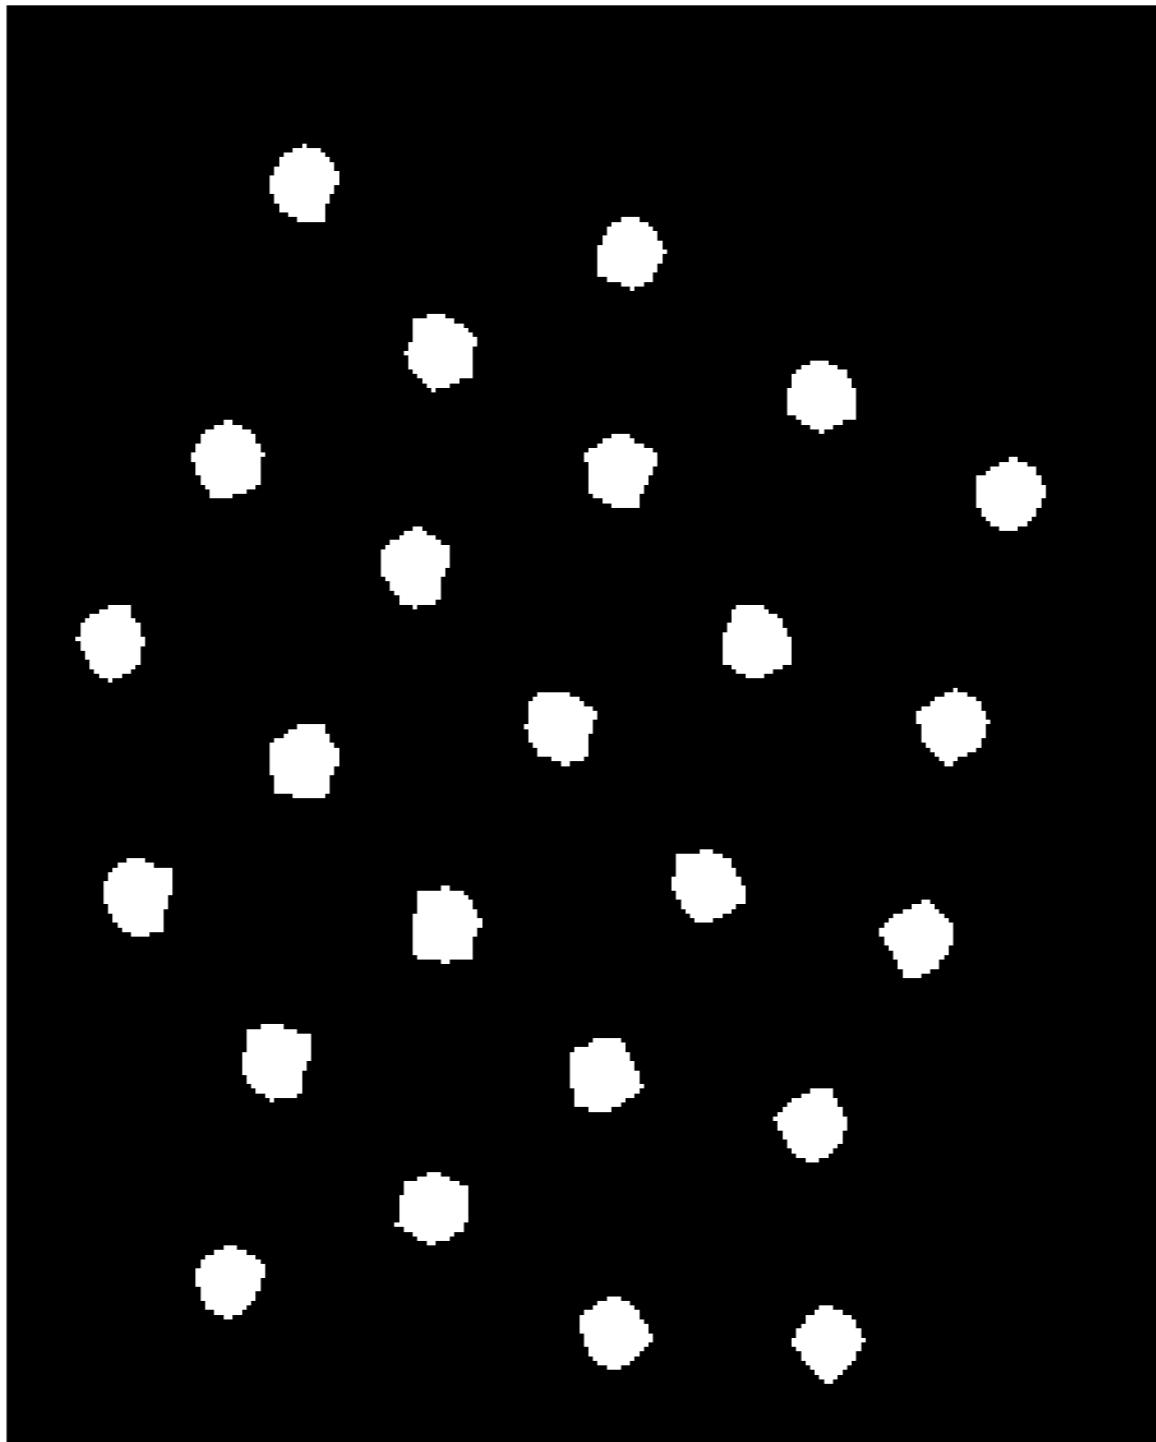


All black pixel are definitely background pixel (i.e. not coins) —> sure\_bg

# Watershed segmentation

## Otsu thresholding + opening + distance transform + thresholding

Distance transform = derived representation of a binary image, where the value of each pixel is replaced by its distance to the nearest background pixel



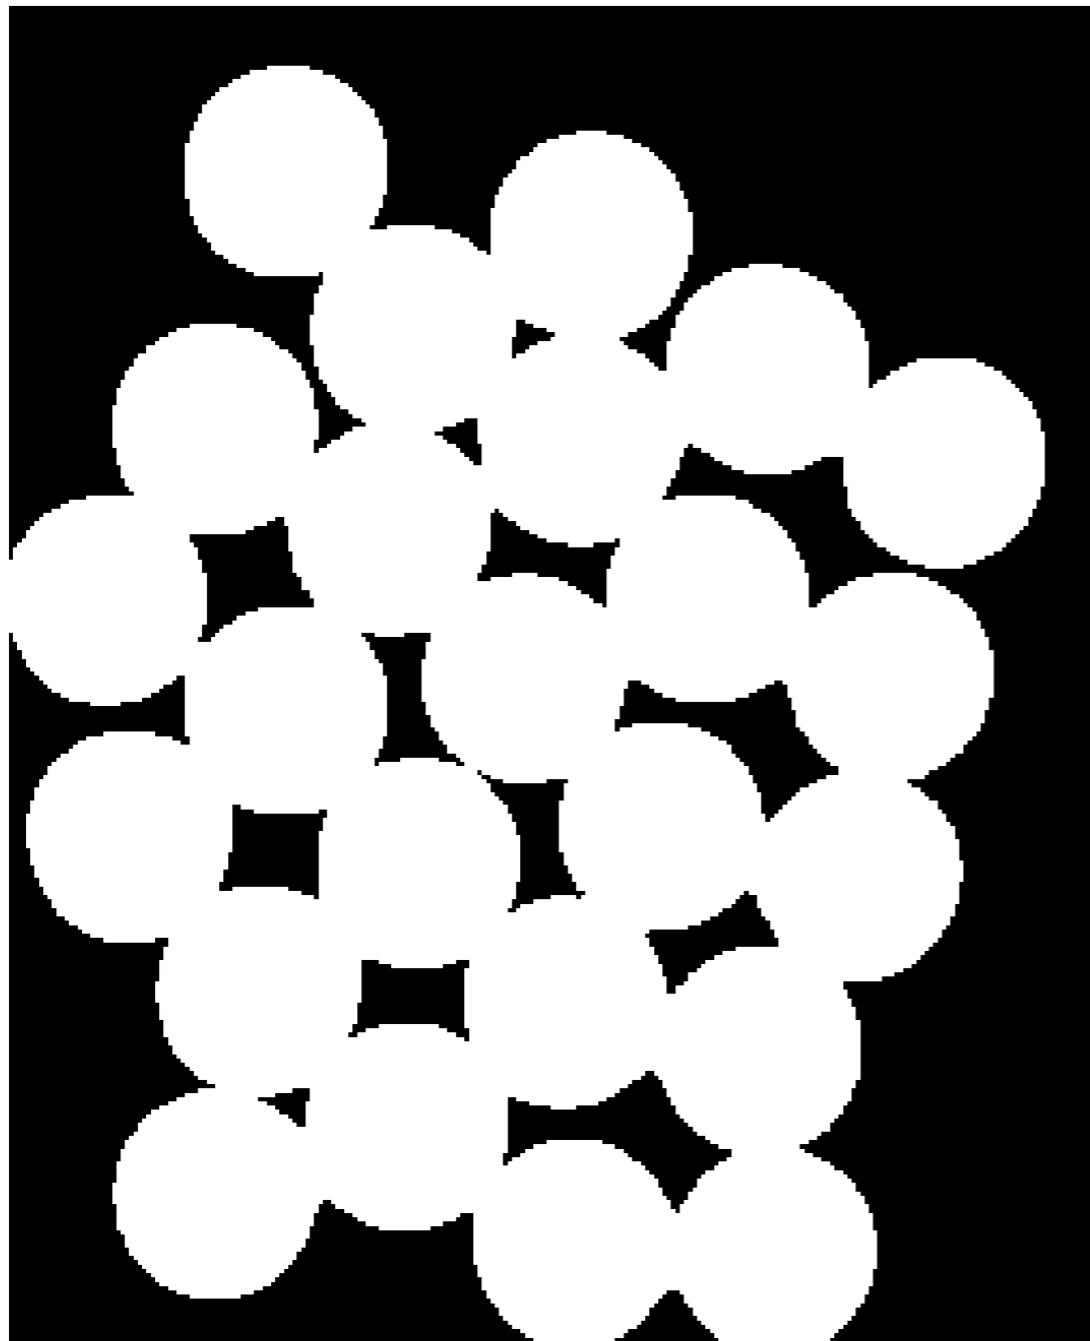
All white pixel are definitely foreground pixel (i.e. coins) —> sure\_fg

# Watershed segmentation

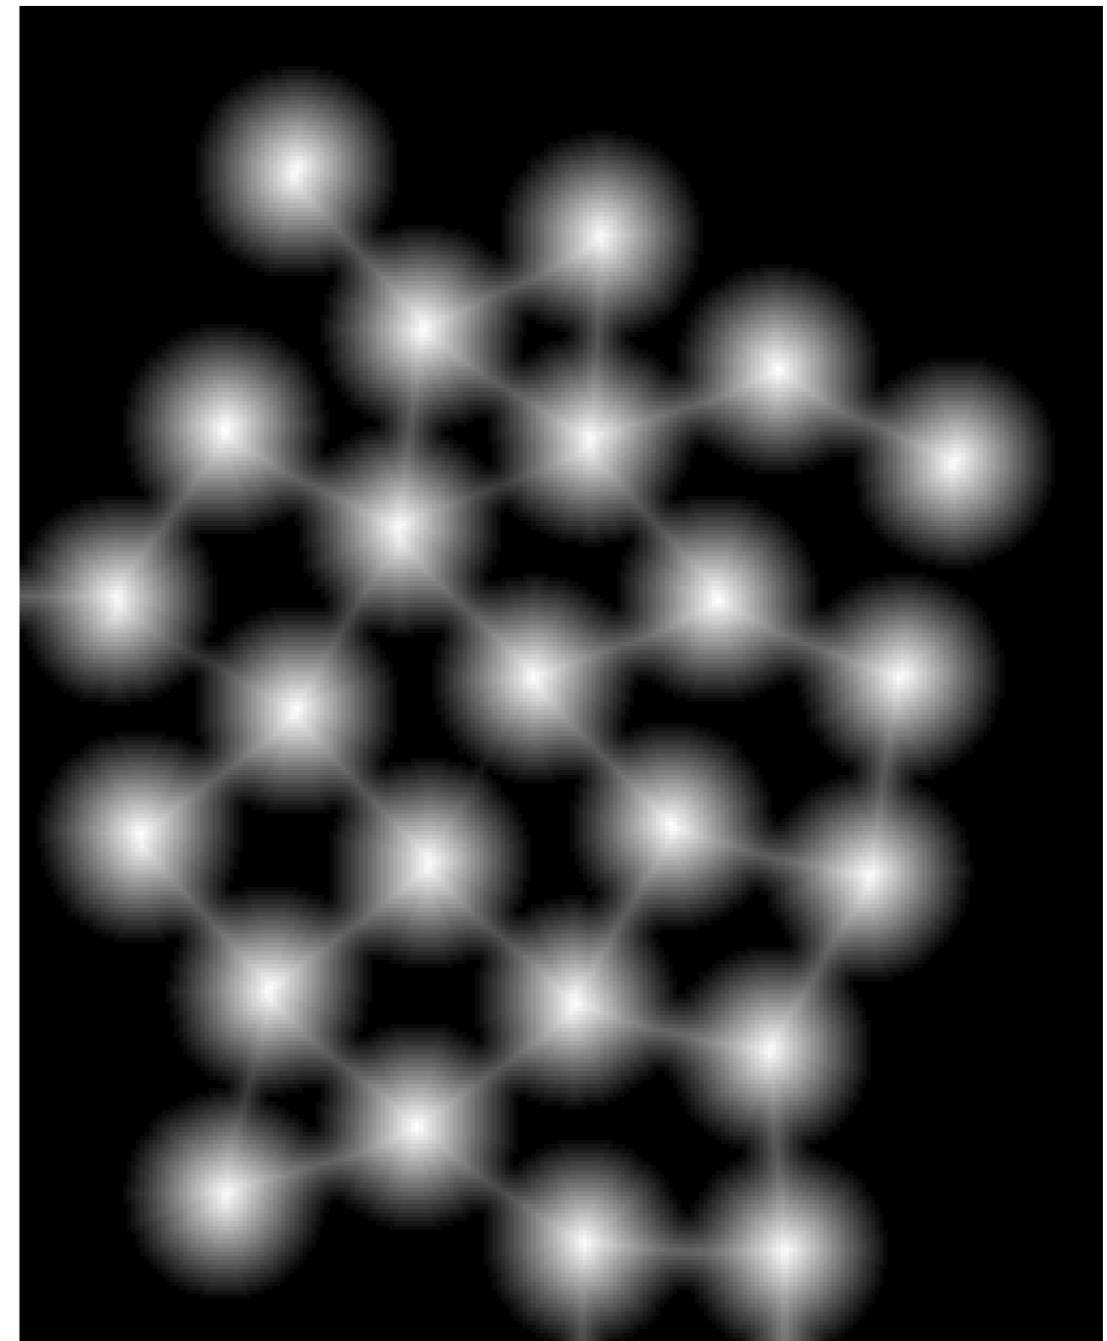
## Distance transform

Distance transform = derived representation of a binary image, where the value of each pixel is replaced by its distance to the nearest background pixel

Opening

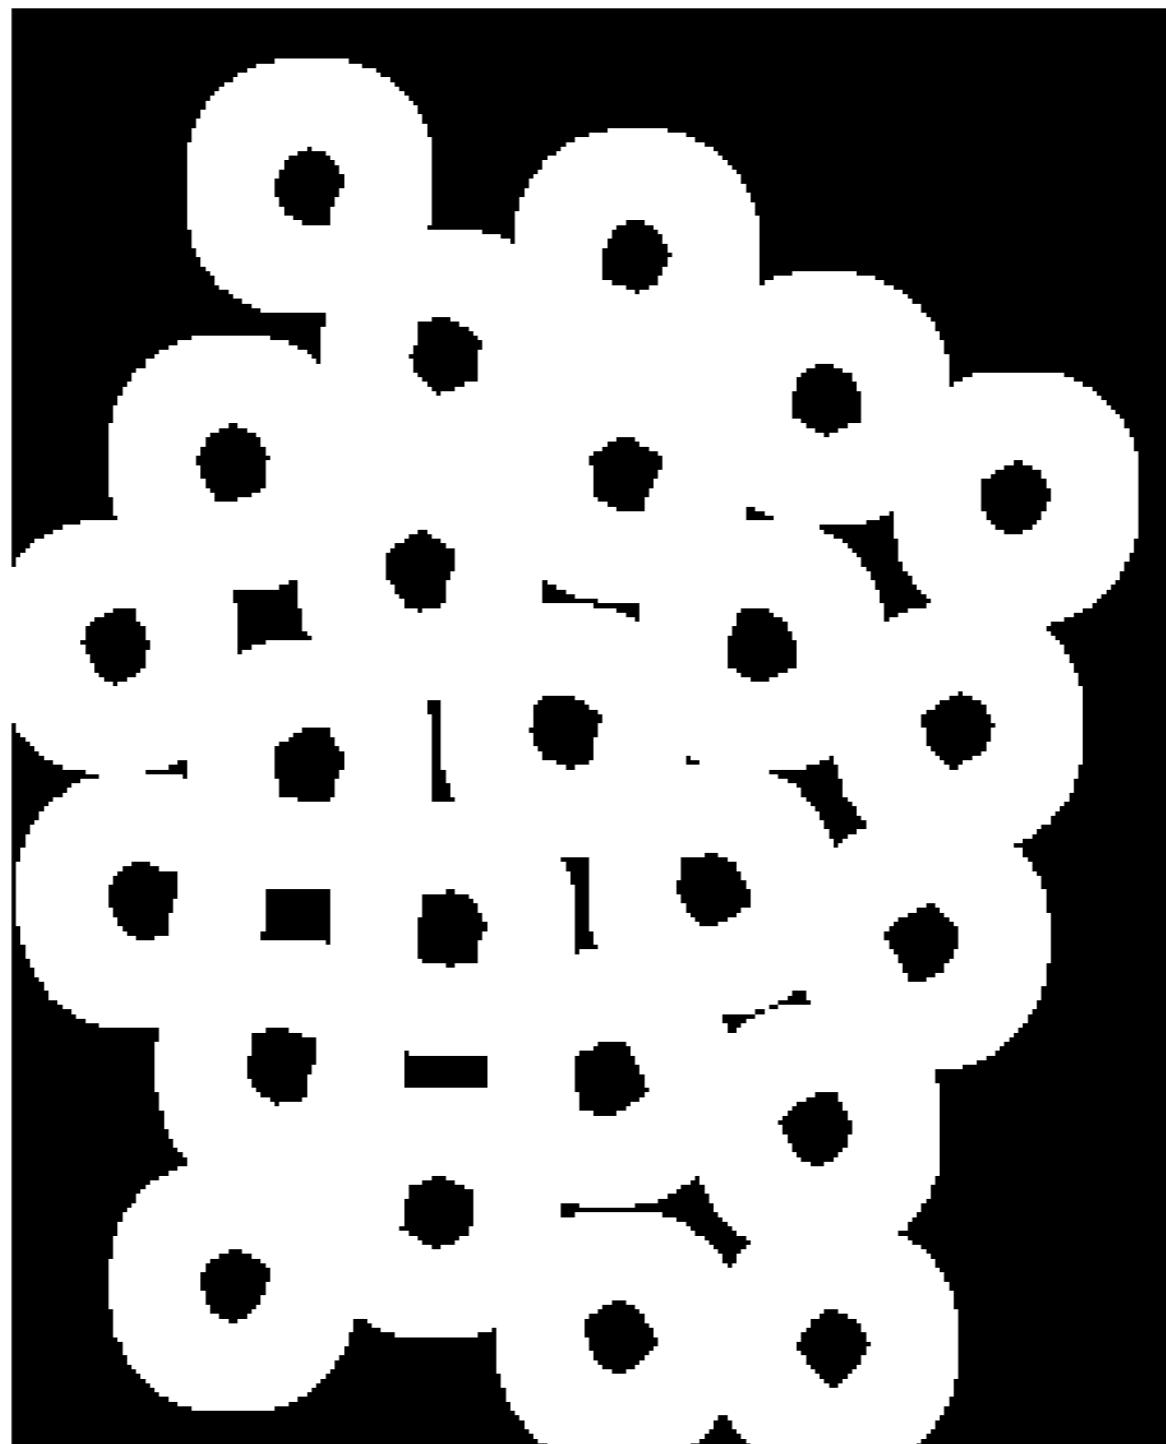


Distance transform



# Watershed segmentation

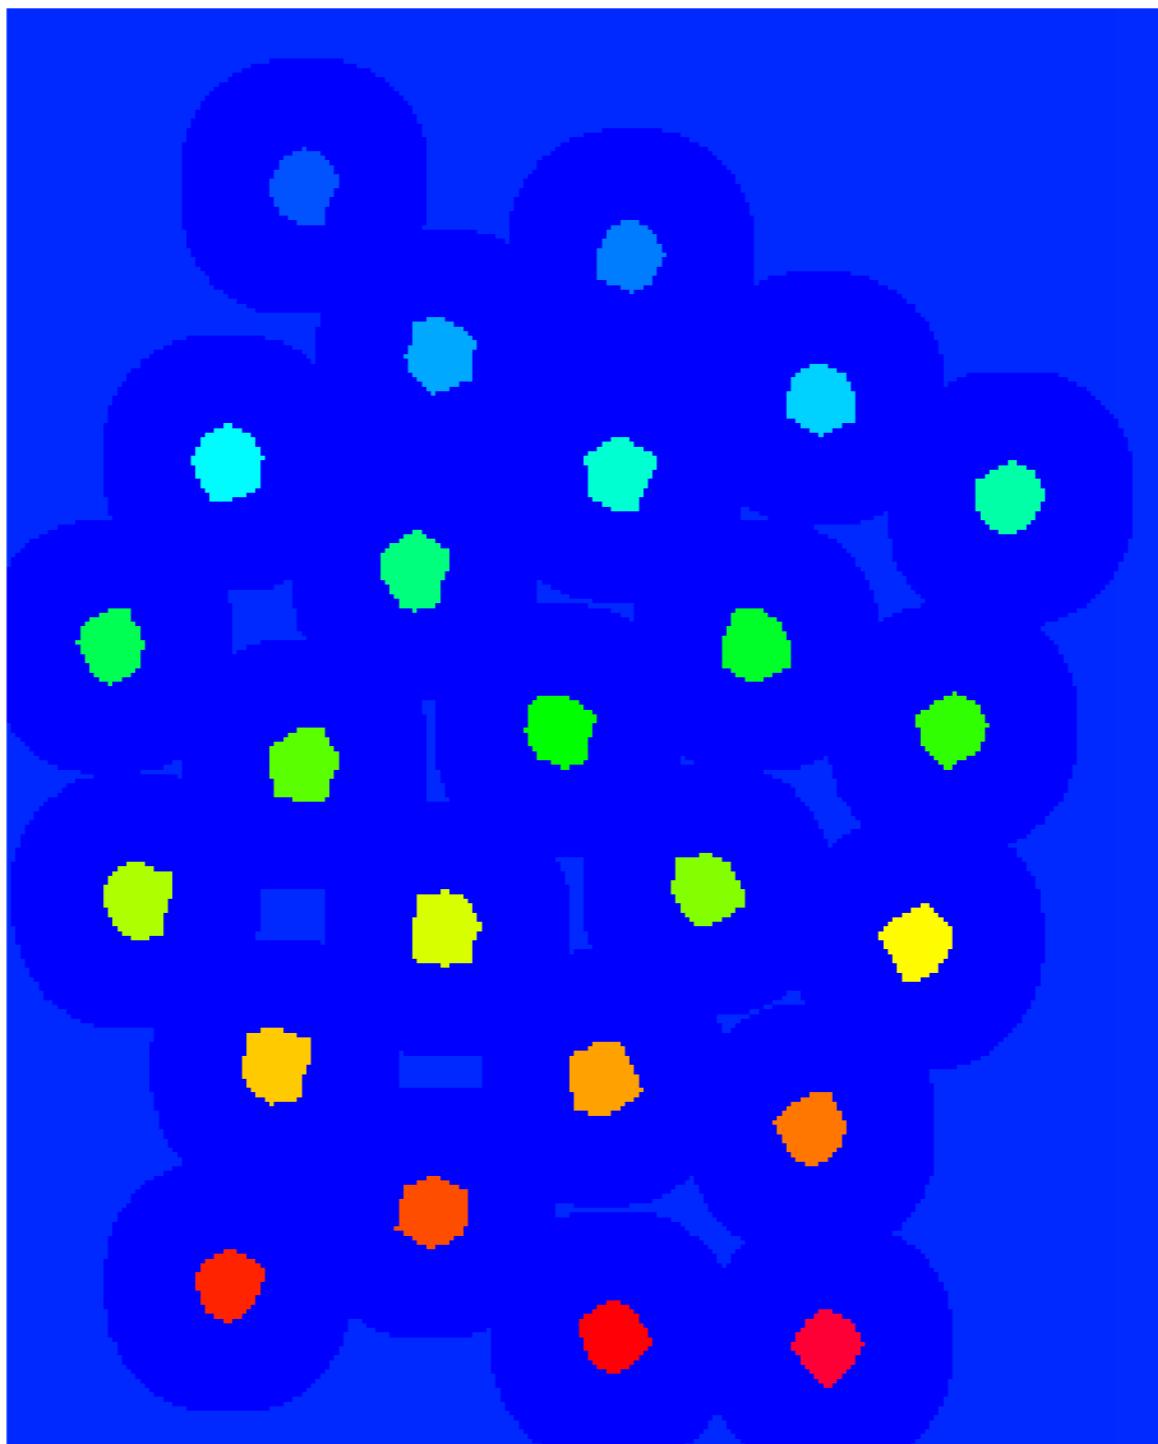
Unknown = sure\_bg-sure\_fg



All white pixel are unknown, we must decide for each pixel

# Watershed segmentation

Markers: one color for each connected region



# Watershed segmentation

Final step: flooding each region to find ridge line

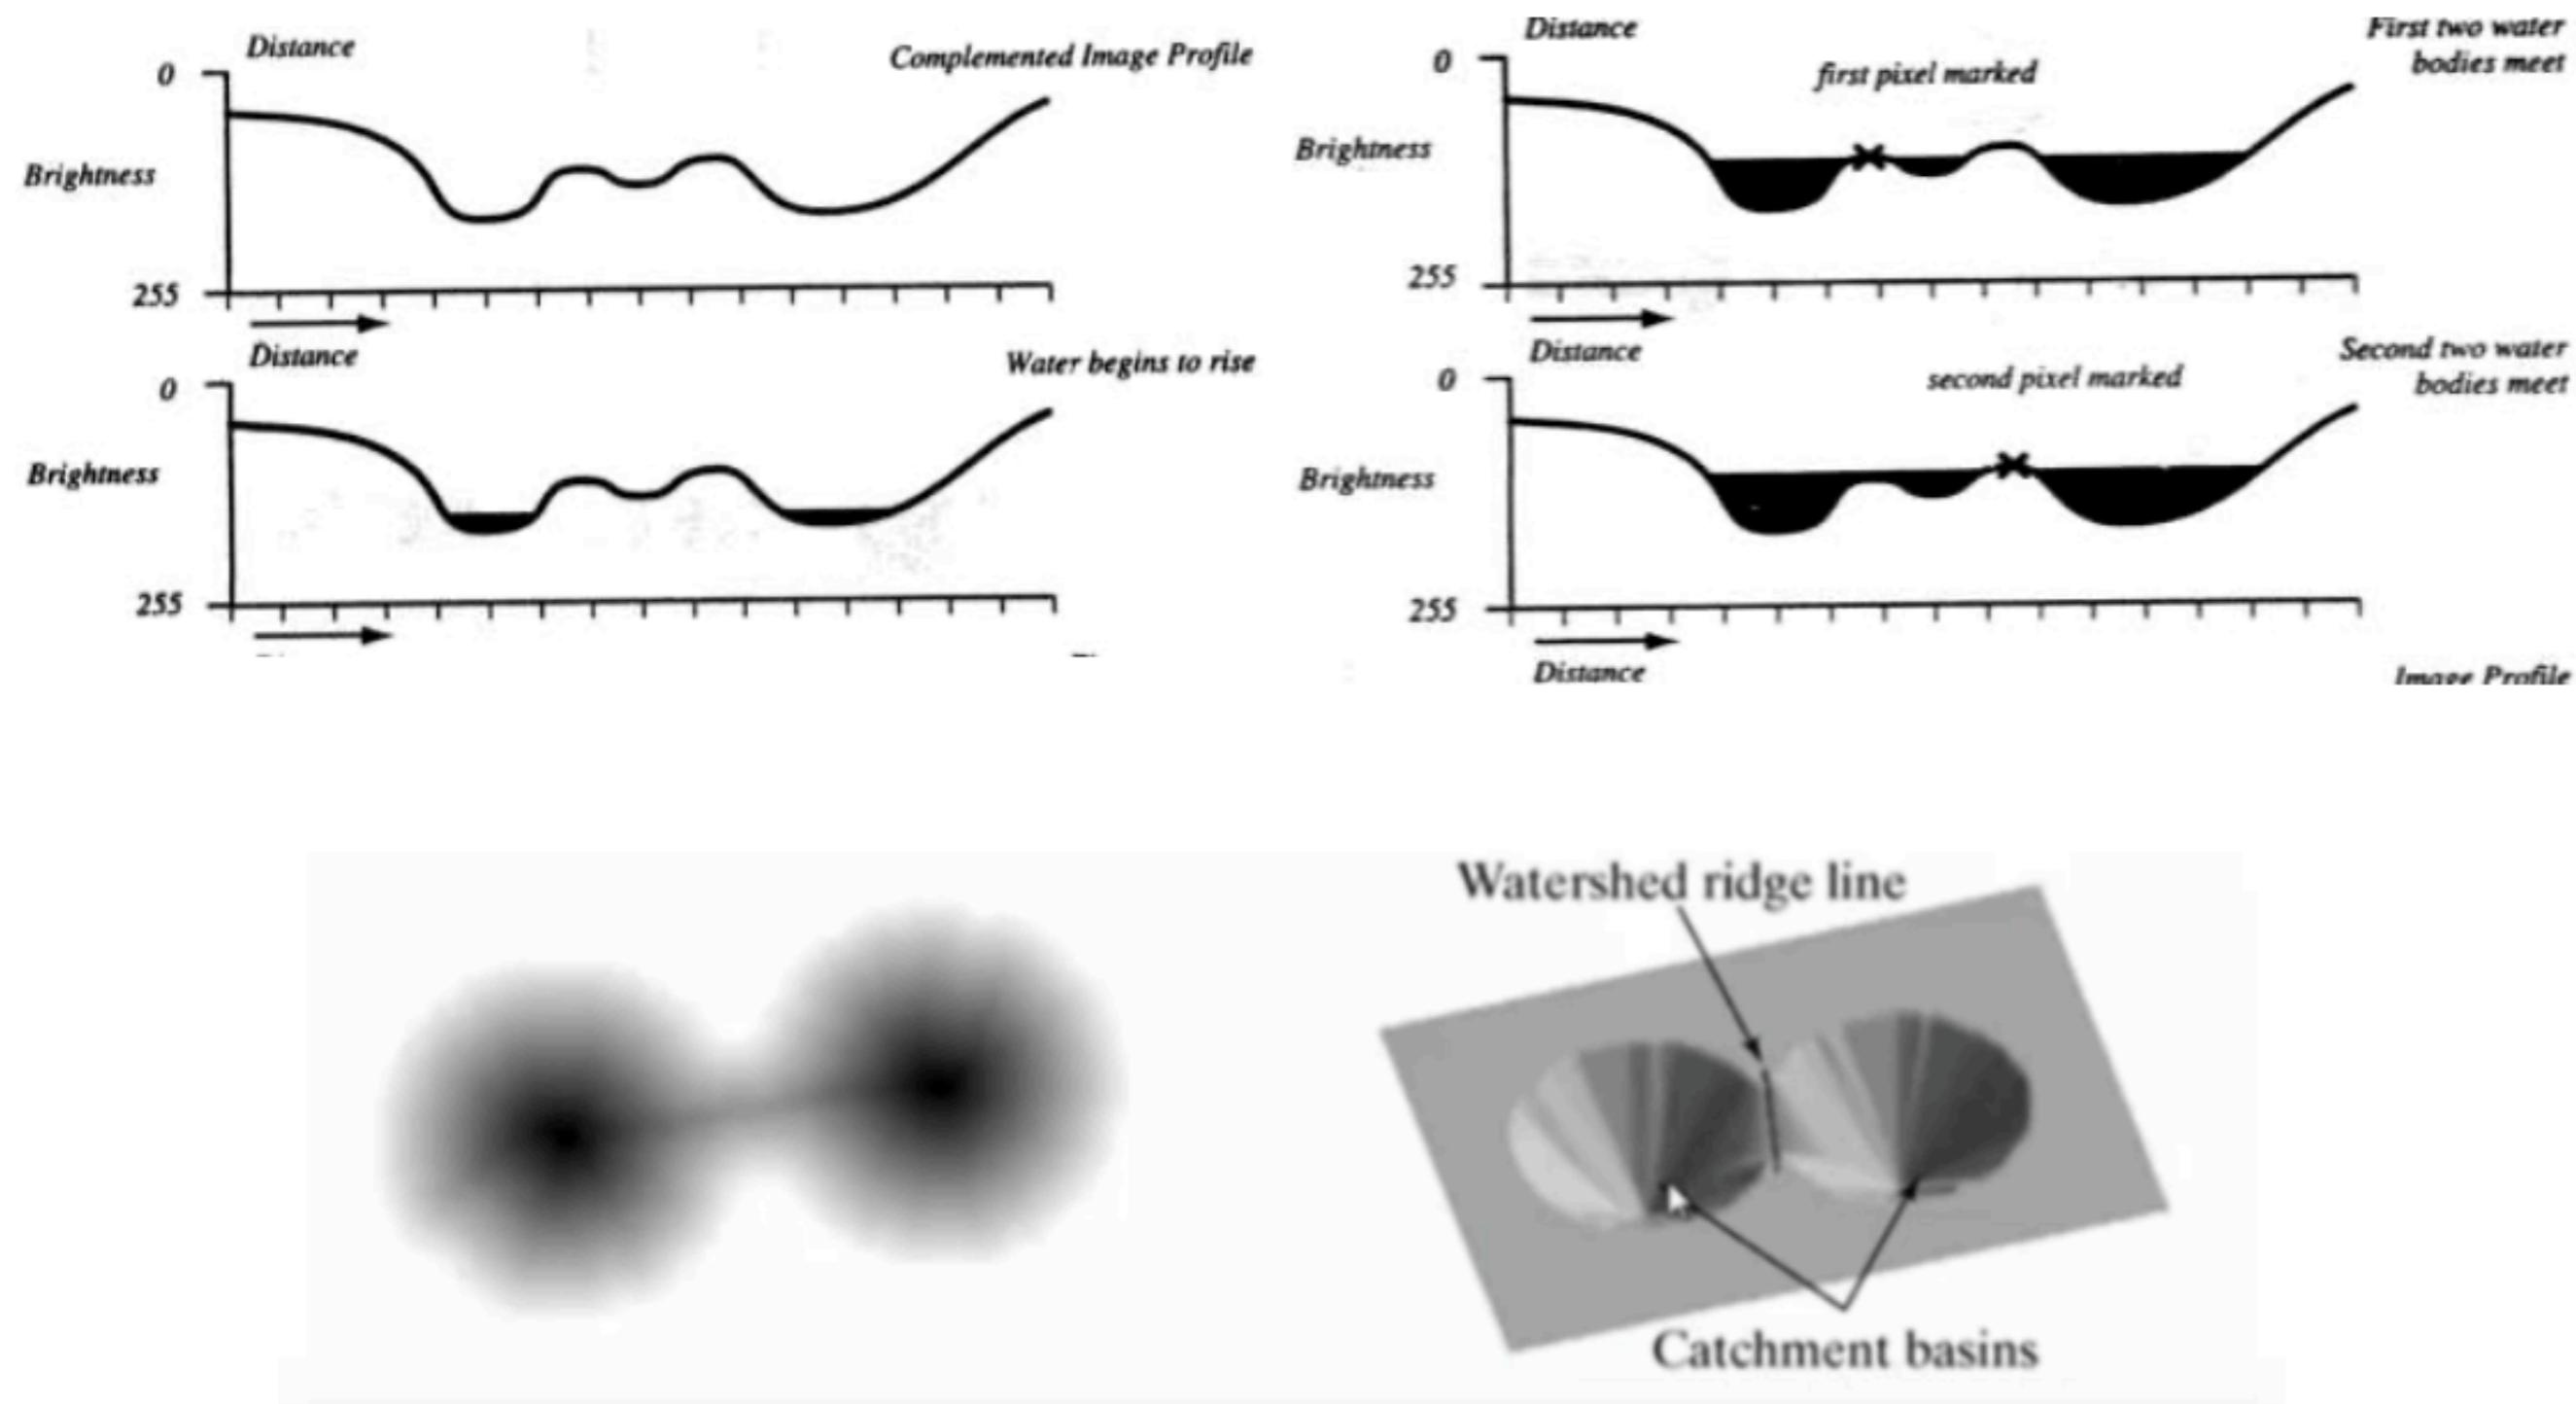


Figure 2. Visualising the watershed: the image on the left can be topographically represented as the image on the right. Source: (Agarwal, 2015)



# **Demo watershed**

## Third segmentation

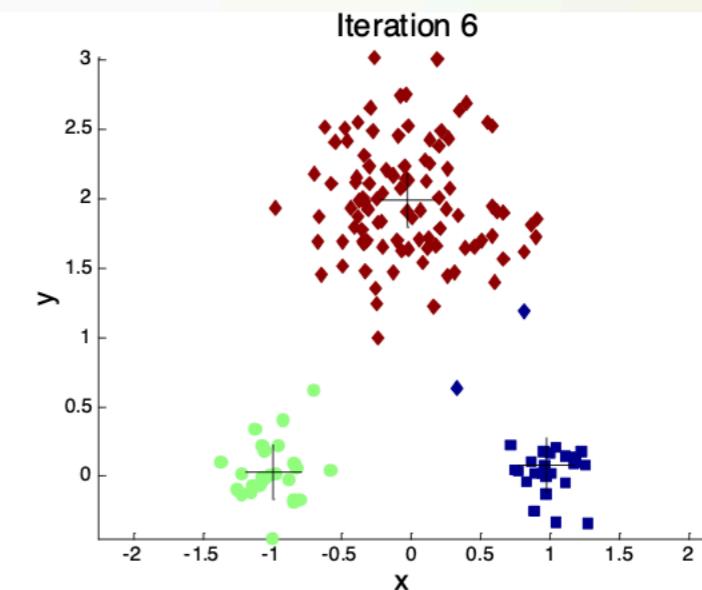
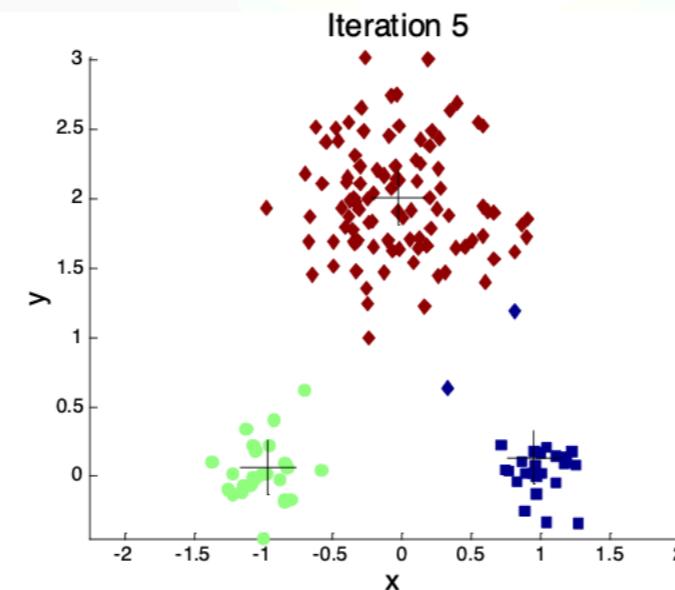
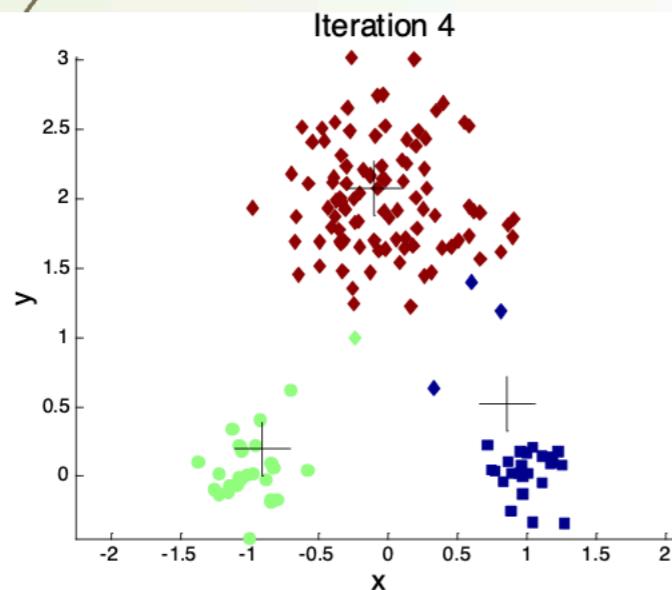
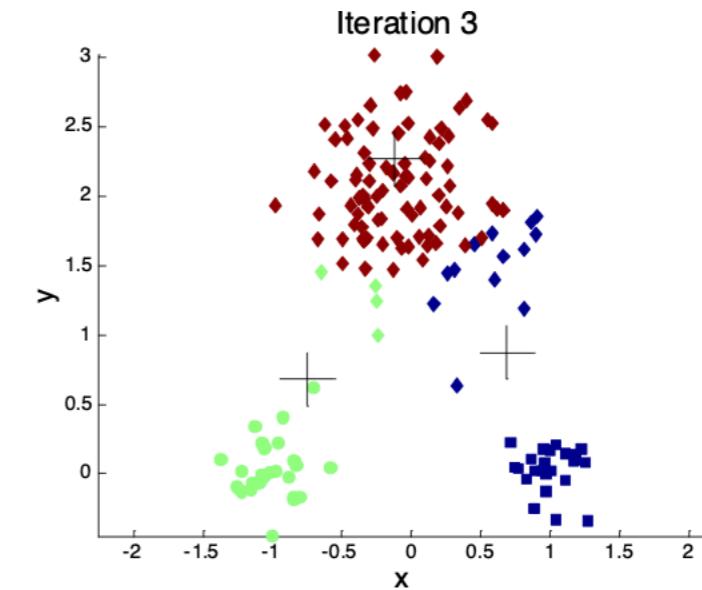
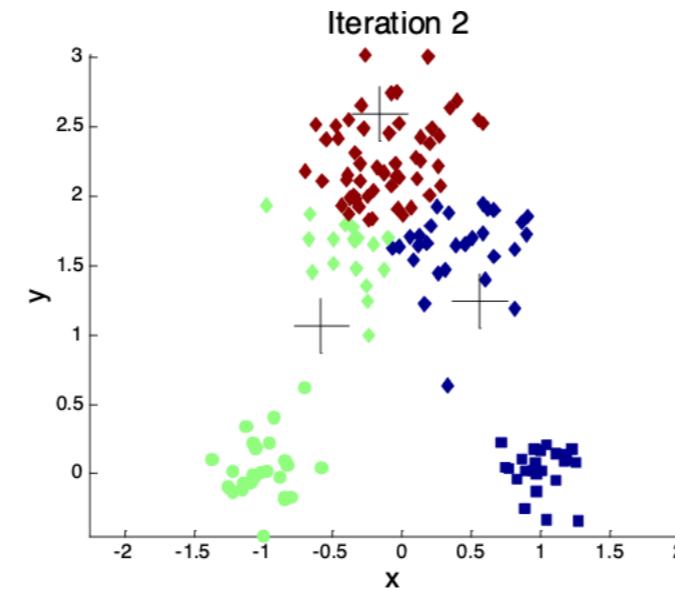
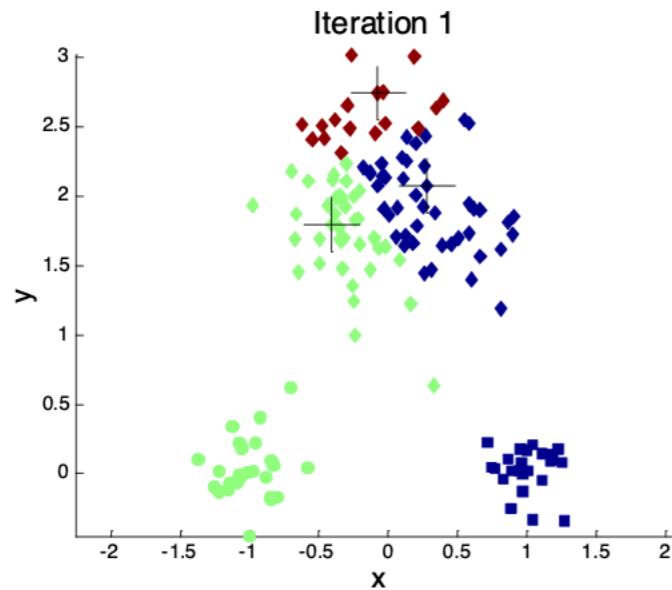
Watershed



# Clustering segmentation

## What is clustering (K-means) ?

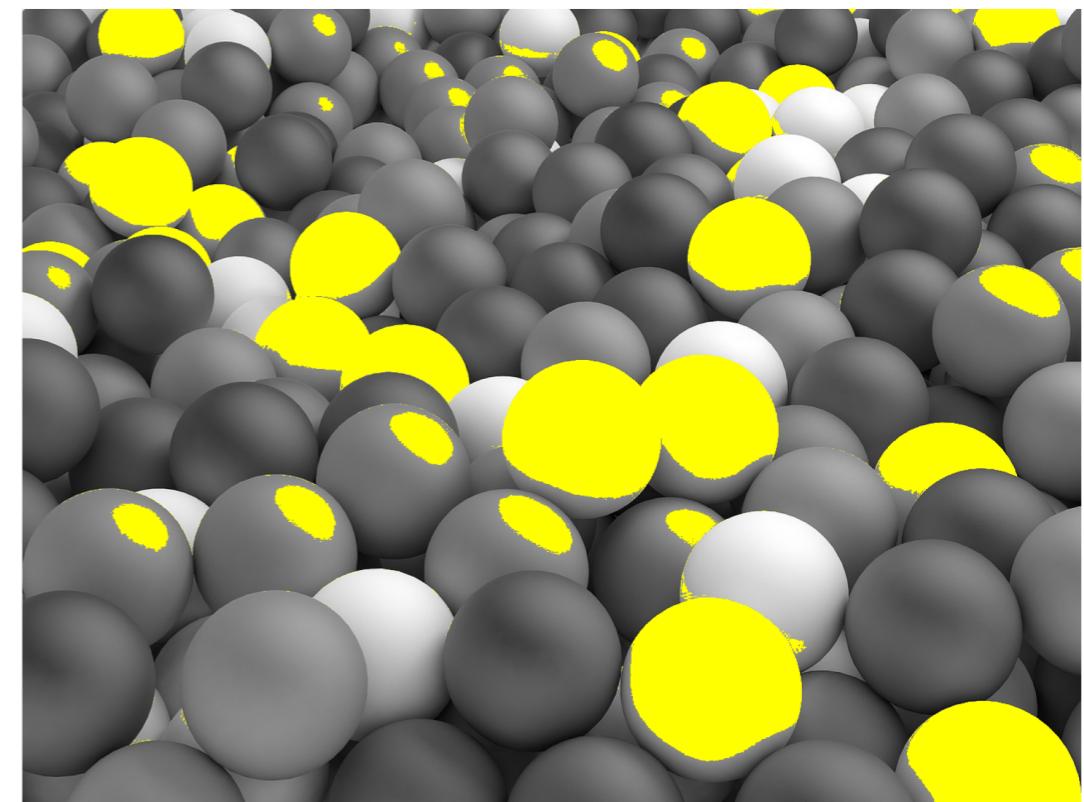
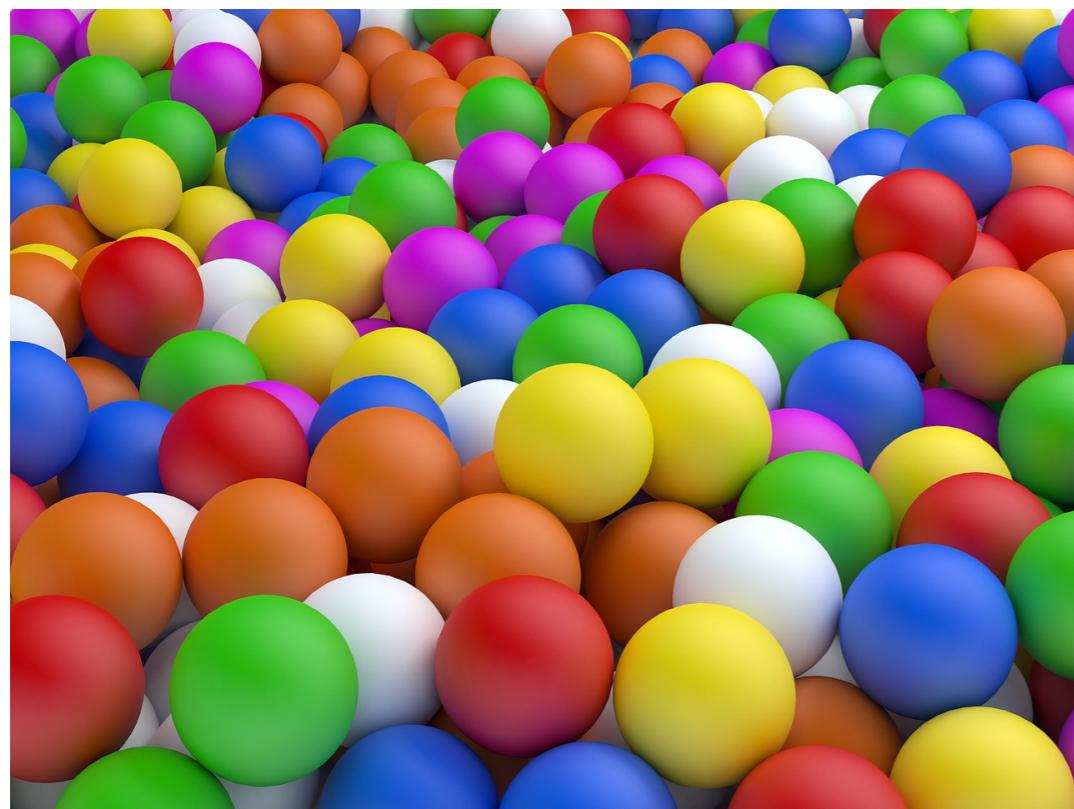
- 1: Select  $K$  points as the initial centroids.
- 2: **repeat**
- 3:   Form  $K$  clusters by assigning all points to the closest centroid.
- 4:   Recompute the centroid of each cluster.
- 5: **until** The centroids don't change



# Clustering segmentation

**K-means on color image → color quantization, i.e. reduce number of colors**

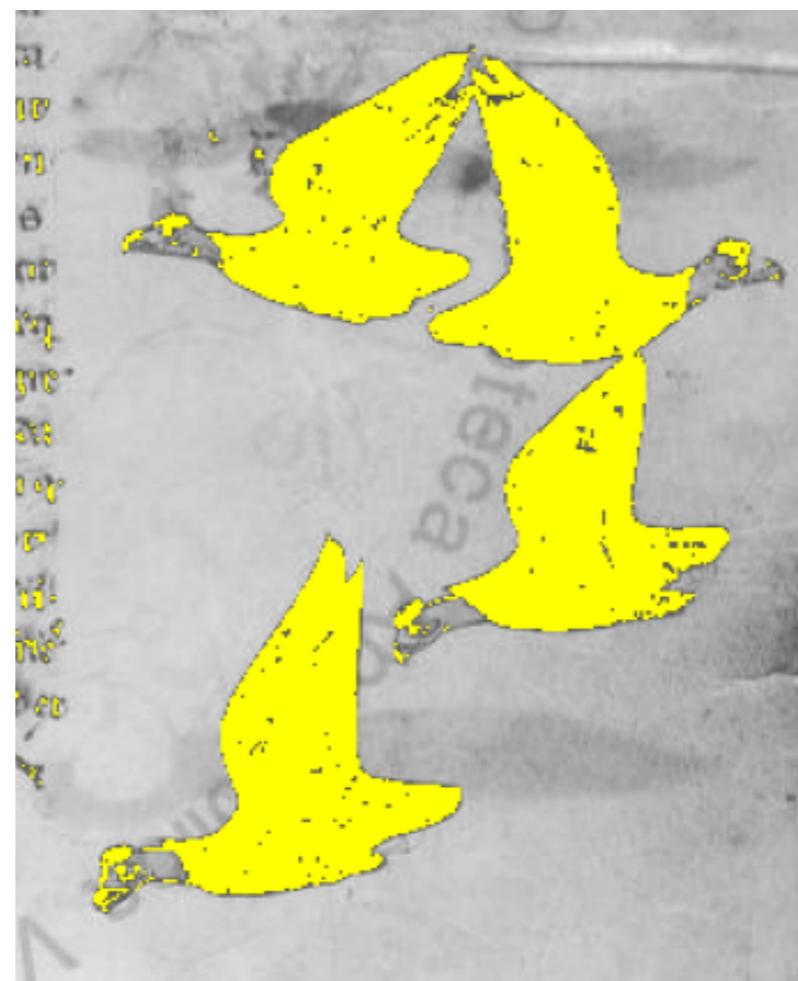
**Example: K=7, display only the class that corresponds to yellow**

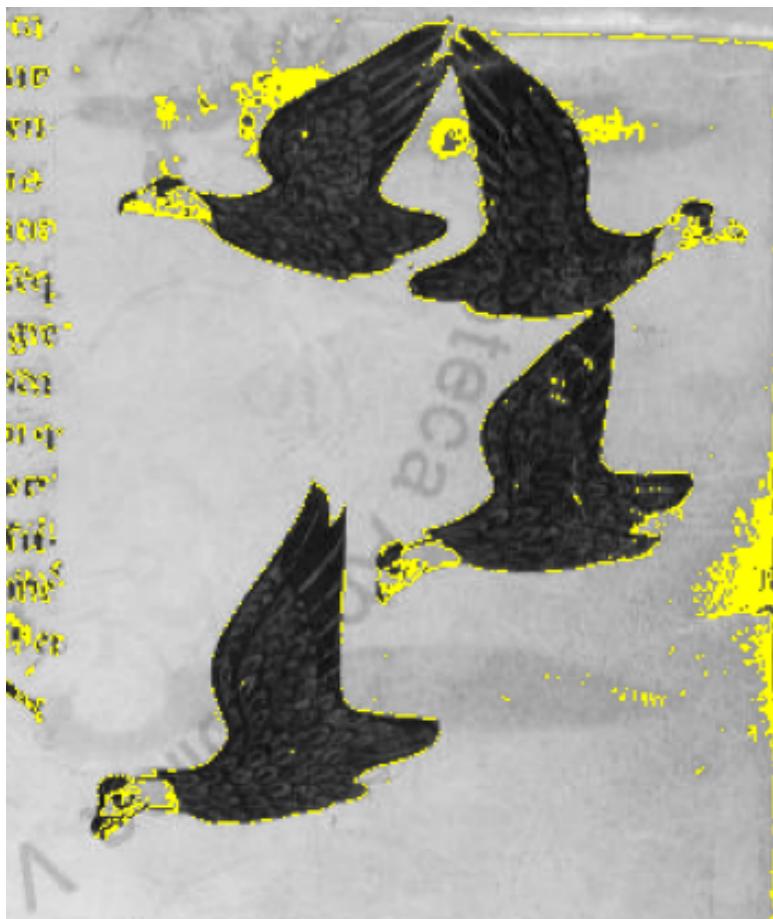




## Fourth segmentation

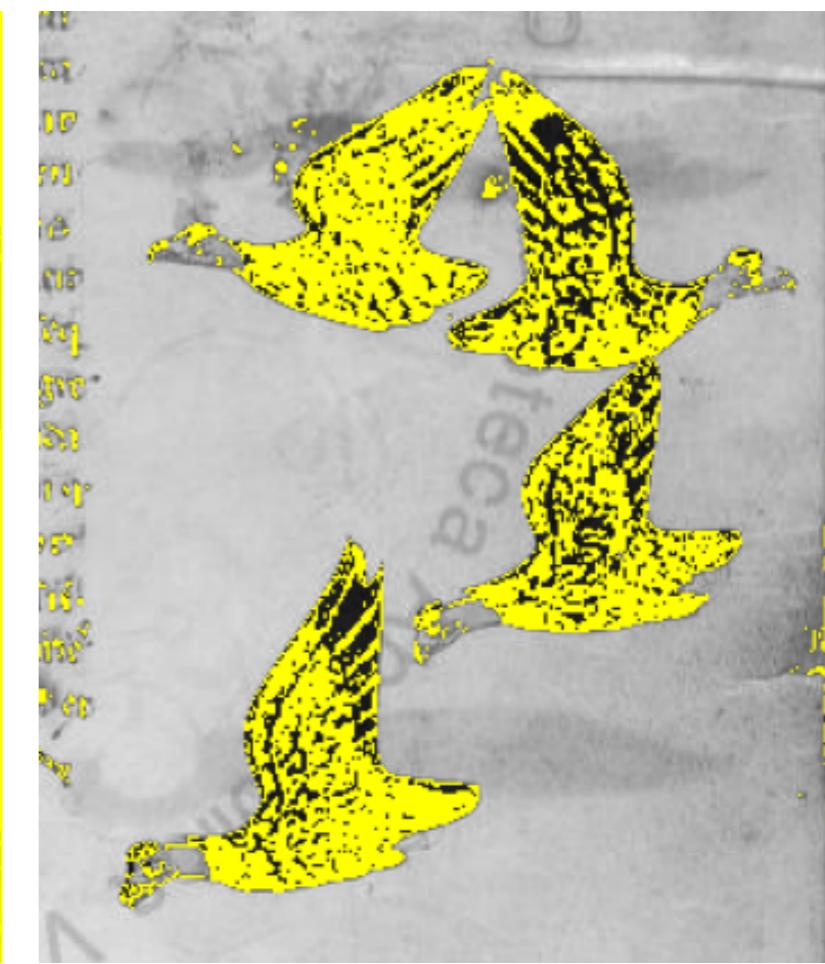
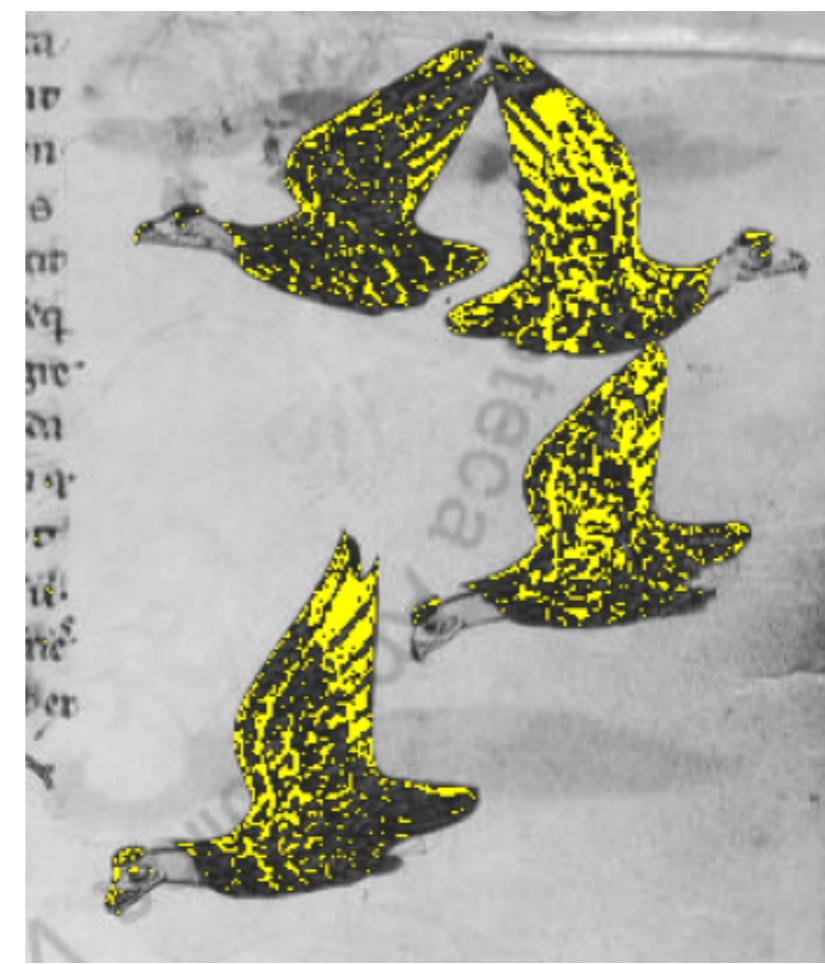
Clustering (K=3)





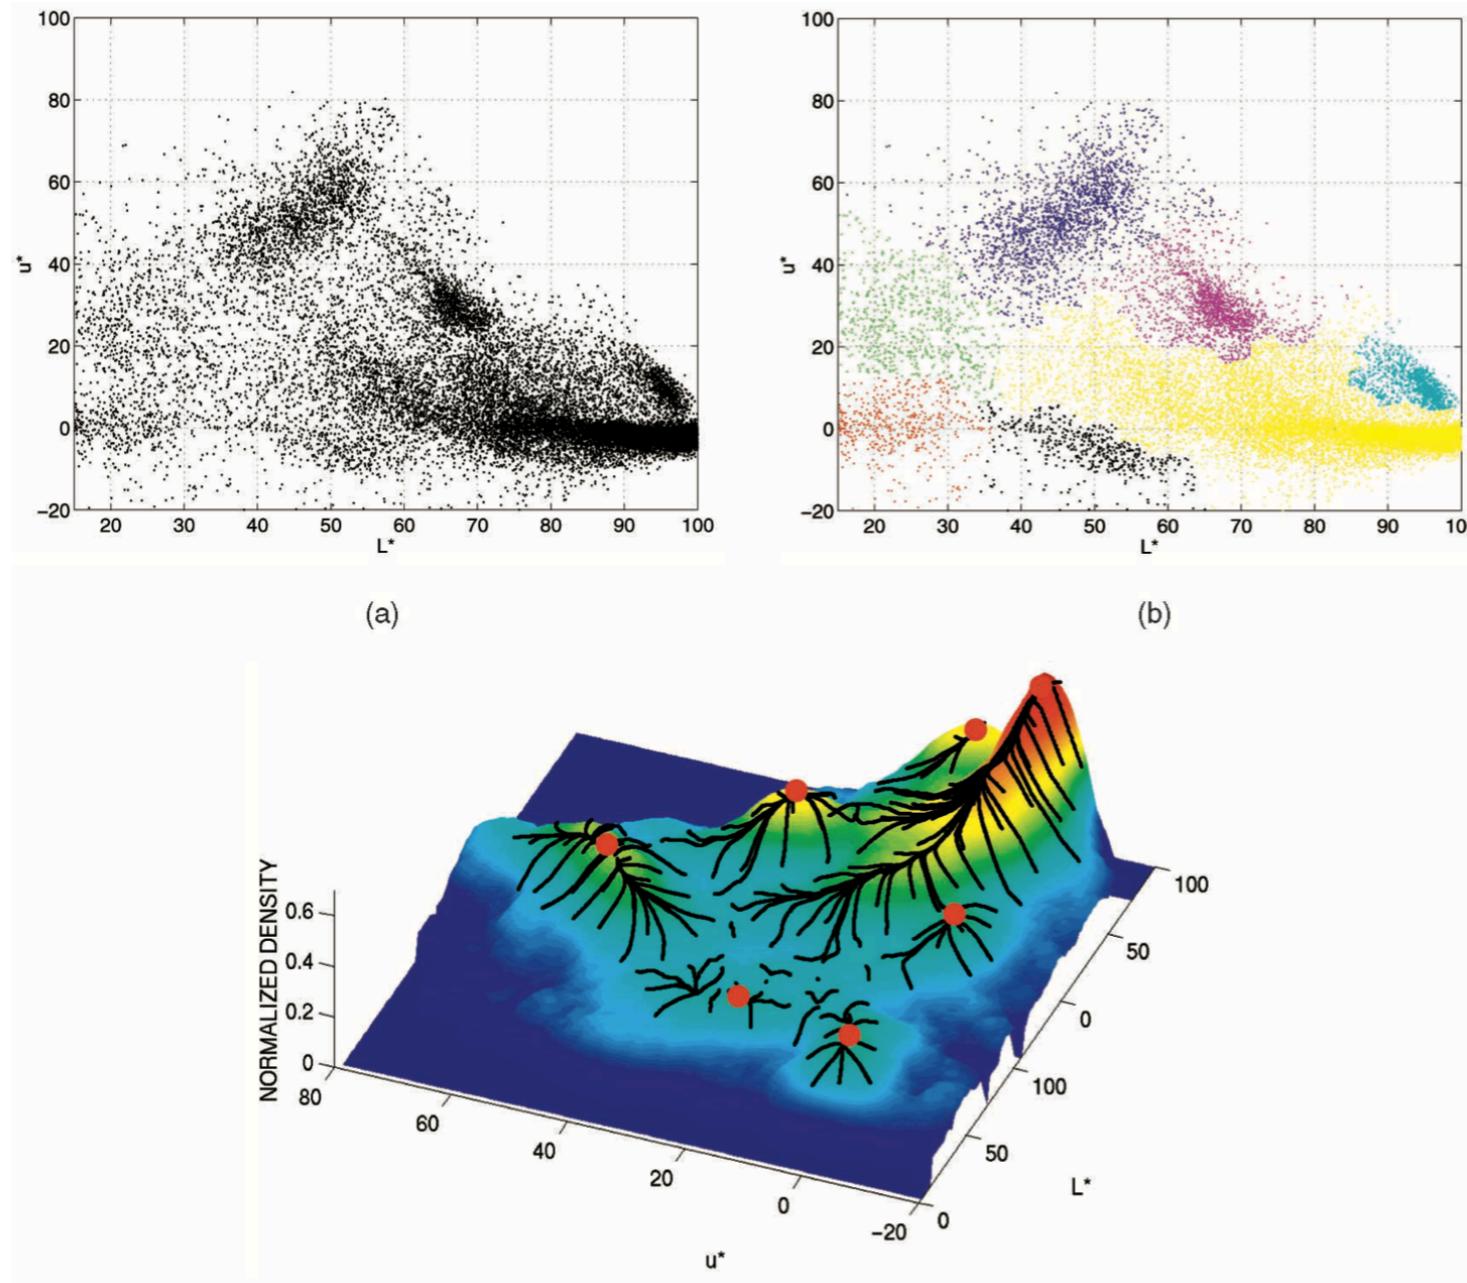
## Fourth segmentation

Clustering (K=5)



# Mean Shift segmentation

Variante non-parametrica al clustering

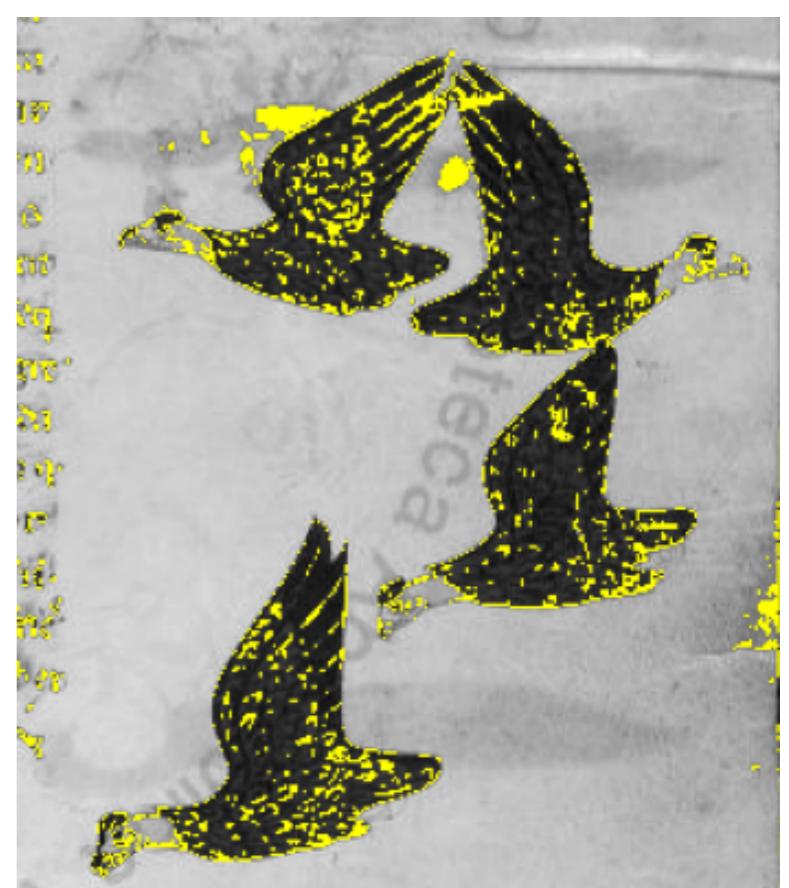
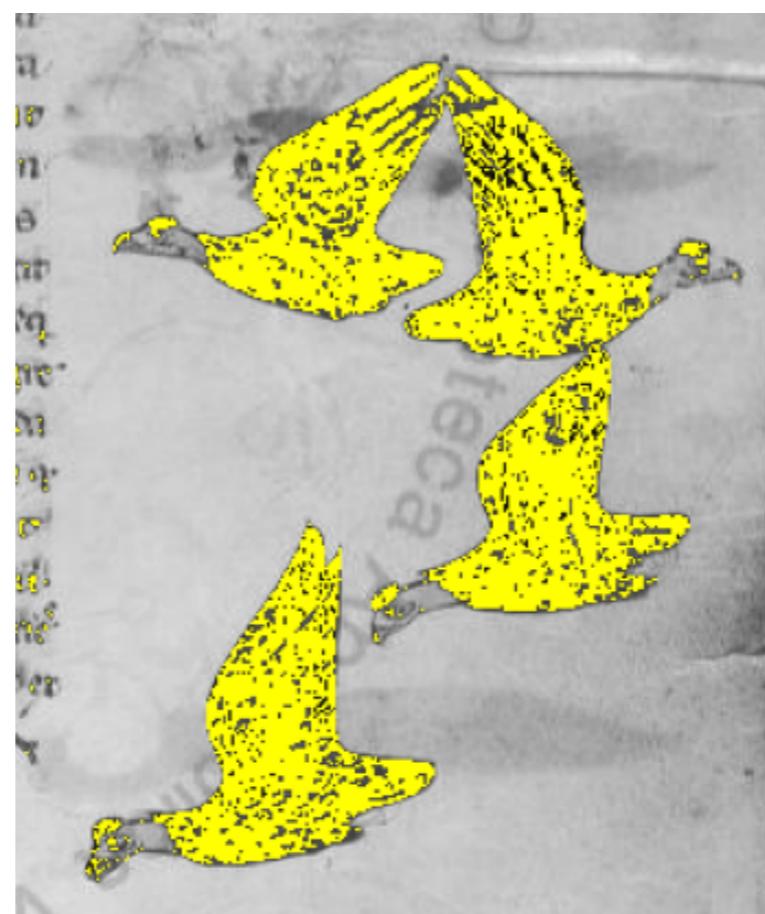
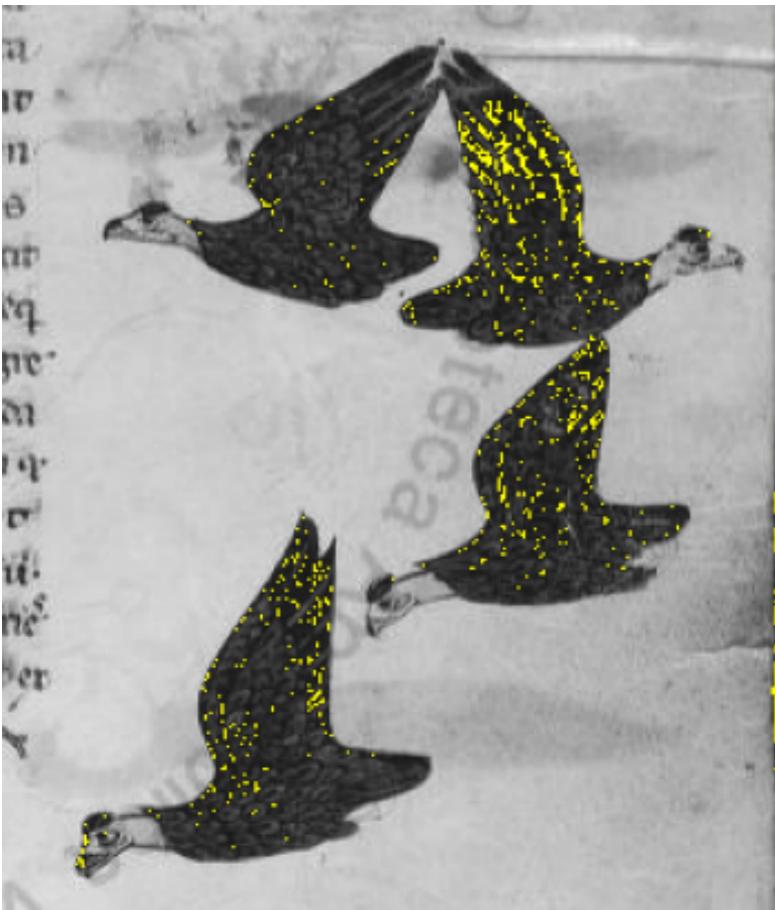


No need to pre-define the number of clusters



## Fifth segmentation

Mean Shift (number of clusters=9)



# **Demo clustering+mean shift**